

# CSC265 F18: Assignment 5

Due: November 21, by midnight

## Guidelines: (read fully!!)

- Your assignment solution must be submitted as a *typed* PDF document. Scanned handwritten solutions, solutions in any other format, or unreadable solutions will **not** be accepted or marked. You are encouraged to learn the L<sup>A</sup>T<sub>E</sub>X typesetting system and use it to type your solution. See the course website for L<sup>A</sup>T<sub>E</sub>X resources. Solutions typed using L<sup>A</sup>T<sub>E</sub>X receive 2 bonus marks.
- Your submission should be no more than 6 pages long, in a single-column US Letter or A4 page format, using at least 9 pt font and 1 inch margins.
- To submit this assignment, use the MarkUs system, at URL <https://markus.teach.cs.toronto.edu/csc265-2018-09>
- This is a *group assignment*. This means that you can work on this assignment with *at most one other* student. You are *strongly encouraged* to work with a partner. Both partners in the group should work on and arrive at the solution together. Both partners receive the same mark on this assignment.
- Work on all problems together. For each problem, one of you should write the solution, and one should proof-read and revise it. The first page of your submission must list the *name*, *student ID*, and *UTOR email address* of both group members. It should also list, for each problem, which group member wrote the problem, and which group member proof-read and revised it.
- You **may not** consult any other resources except: your partner; your class notes; your textbook and assigned readings. *Consulting any other resource, or collaborating with students other than your group partner, is a violation of the academic integrity policy!*
- You may use any data structure, algorithm, or theorem previously studied in class, or in one of the prerequisites of this course, by just referring to it, and without describing it. This includes any data structure, algorithm, or theorem we covered in lecture, in a tutorial, or in any of the assigned readings. Be sure to give a *precise reference* for the data structure/algorithm/result you are using.
- Unless stated otherwise, you should justify all your answers using rigorous arguments. Your solution will be marked based both on its completeness and correctness, and also on the clarity and precision of your explanation.

**Question 1.** (15 marks)

Recall that during an  $\text{AVL-INSERT}(x)$  operation in an AVL tree, we first insert the key  $x$  into the tree as we do in any binary search tree (BST). Then we go up the path from the new node to the root and update the balance factor fields until we either reach a node which becomes perfectly balanced after the insertion, or we have to do a single or a double rotation. (Review the notes on AVL trees.)

This second part of an  $\text{AVL-INSERT}(x)$  – updating balance factors and doing rotations – can, in the worst case, double the time it takes to insert a new node. Your goal in this question is to show that this is not true in an amortized sense.

Use the potential function method to show that in any sequence of  $m$  consecutive  $\text{AVL-INSERT}(x)$  operations, starting from an empty AVL tree, the total number of balance factor updates is bounded by  $O(m)$ . I.e. if  $t_i$  is the number of nodes whose balance factor is changed during the  $i$ -th consecutive insert (not counting the new node), show that  $\sum_{i=1}^m t_i = O(m)$ . Be precise about the potential you are using. Show that the potential is 0 for an empty AVL tree, and is never negative. Derive the claimed bound on  $\sum_{i=1}^m t_i$  by analyzing the potential differences.

**Question 2.** (20 marks)

We define a simple mergeable heap data structure, called the TREE-HEAP. Unlike Binomial Heaps, a TREE-HEAP is implemented as a single (but *not necessarily binary*) tree. A TREE-HEAP is a tree which stores the keys of the heap elements at the nodes, so that the *min-heap property* is satisfied: the key stored at every node is at most the keys of its children. This is the only condition the TREE-HEAP needs to satisfy!

We will discuss three operations on TREE-HEAP's:

- $\text{INSERT}(T, x)$ , which inserts a new item with key  $x$ ;
- $\text{UNION}(T_1, T_2)$  which merges the heaps  $T_1$  and  $T_2$  into a new TREE-HEAP  $T$ ;
- $\text{EXTRACT-MIN}(T)$  which returns the heap element with the smallest key in  $T$ , and removes it from  $T$ .

As usual, we implement  $\text{INSERT}(T, x)$  by creating a new heap  $T'$  with a single node storing the key  $x$ , and calling  $\text{UNION}(T, T')$ . The operation  $\text{UNION}(T_1, T_2)$  is very simple: we compare the keys of the roots of  $T_1$  and  $T_2$ , and make the root with the larger key a child of the root with the smaller key. You can assume that when  $T_1$  becomes a subtree of  $T_2$ , it is placed as the leftmost subtree of the root, and vice versa. See Figure 1 for an example.




Figure 1: An example of a union of two TREE-HEAP's.

Since a TREE-HEAP is not a binary tree, we implement it using the leftmost child – right sibling representation, similarly to how binomial trees were implemented. In particular, each node  $u$  of a TREE-HEAP  $T$  has the following fields:

- $u.key$  stores the key of  $u$ ;

- $u.lchild$  stores a pointer to the leftmost child of  $u$ , if  $u$  has any children, and NIL otherwise;
- $u.rsibl$  stores a pointer to the sibling of  $u$  directly to its right, if  $u$  has any siblings to the right, and NIL otherwise.

You should make sure that with this representation you know how to implement  $\text{UNION}(T_1, T_2)$  by changing a constant number of pointers.

The  $\text{EXTRACT-MIN}(T)$  operation is the trickiest one, and can be very slow if not implemented carefully. Below we explore two options.

**Part a.**

Suppose we implement  $\text{EXTRACT-MIN}(T)$  by removing the root of  $T$  (and storing its key to return once we are done); then we are left with the subtrees of the root of  $T$ : let's call them  $T_1, \dots, T_d$ , listed from left to right. Here  $d$  denotes the degree of the root of  $T$  before it was deleted. Now we successively perform  $T'_2 = \text{UNION}(T_1, T_2)$ ;  $T'_3 = \text{UNION}(T'_2, T_3)$ ,  $\dots$ ,  $T'_d = \text{UNION}(T'_{d-1}, T_d)$ , and return  $T'_d$  as the new heap.

For every  $m$ , show a sequence of  $m$  INSERT and EXTRACT-MIN operations, starting from an empty heap, so that the total time complexity of executing the sequence using the algorithms above is  $\Omega(m^2)$ . I.e. show that the amortized complexity is  $\Omega(m)$  per operation. Justify your answer, and be specific about the exact sequence of operations.

**Part b.**

Since the algorithm for  $\text{EXTRACT-MIN}(T)$  above turned out to be too slow, let us try another one. Again we remove the root of  $T$  (and store its key to return once we are done); then we are left with the subtrees of the root of  $T$  that we removed, denoted  $T_1, \dots, T_d$ , which we need to merge into a single tree. We will now do this in two stage. First we merge them in pairs: we run  $T'_1 = \text{UNION}(T_1, T_2)$ ,  $T'_2 = \text{UNION}(T_3, T_4)$ ,  $\dots$ ,  $T'_k = \text{UNION}(T_{2k-1}, T_{2k})$ , where  $k = \lfloor d/2 \rfloor$ . If  $d$  is odd, then we set  $T'_{k+1} = T_d$ . This completes the first stage. In the second stage we merge the  $T'_i$  trees as we did before: we set  $T''_2 = \text{UNION}(T'_1, T'_2)$ ;  $T''_3 = \text{UNION}(T''_2, T'_3)$ ,  $\dots$ ,  $T''_\ell = \text{UNION}(T''_{\ell-1}, T'_\ell)$ , where  $\ell = \lceil d/2 \rceil$ . Finally we return  $T''_\ell$ .

Show that in any sequence of  $m$  INSERT and EXTRACT-MIN operations, starting with the empty heap, the amortized cost of an INSERT is  $O(1)$ , and the amortized cost of an EXTRACT-MIN is  $O(\sqrt{m})$ . Prove this using the potential function method. Be precise about the potential you are using. Show that the potential is 0 for an empty TREE-HEAP, and is never negative. Derive the claimed bound on the amortized complexity by analyzing the potential differences.

For the above, you can assume that the actual cost of an  $\text{EXTRACT-MIN}(T)$  equals the number of children of the root of  $T$ , and that the actual cost of  $\text{INSERT}(T, x)$  equals 1.

HINT: Consider assigning every node  $u$  of  $T$  a potential  $\phi(u) = 1 - \min\{d(u), n^\alpha\}$  for some constant  $0 < \alpha \leq 1$ . Here  $d(u)$  equals the number of children of  $u$ , and  $n$  equals the number of nodes in  $T$ . Then define the potential of  $T$  as the sum of the potentials of its nodes. Specify a precise value for  $\alpha$  and prove all the necessary properties of the potential. You can use the inequality  $(a + b)^\alpha \leq a^\alpha + b^\alpha$ , valid for all  $a, b \geq 0$  and  $0 < \alpha \leq 1$ .