

# The Maximum Ratio Subtree Problem

We will use mergeable heaps to solve the following optimization problem defined on trees.

We are given a tree  $T$  with  $n$  nodes, and each node  $u$  of the tree stores two positive numbers,  $a_u$  and  $b_u$ . In general in this problem you don't need the tree to be binary, but to keep things simple, let's assume that it actually is binary.

We will call  $T'$  a subtree of  $T$  if it is connected and contains a (nonempty) subset of the nodes of  $T$  together with the edges between them. In Figure 1 you see some examples of subtrees and one example of a non-subtree.




Figure 1: The first two pictures (from the left) show subtrees (marked in red). The last picture is not a subtree because it is not connected.

For a subtree  $T'$ , let  $r(T') = \frac{\sum_{u \in T'} a_u}{\sum_{u \in T'} b_u}$ . I.e.  $r(T')$  is the ratio of the sum of the  $a$  values to the sum of the  $b$  values, where both sums are over the nodes of  $T'$ .

For any node  $u$  of  $T$ , let  $\Gamma(u)$  be the set of all subtrees of  $T$  whose root is  $u$ . Let  $u^*$  be the root of  $T$ . Our goal is to find the value

$$r(u^*) = \max\{r(T') : T' \in \Gamma(u^*)\}.$$

In fact, we will get an algorithm which computes

$$r(u) = \max\{r(T') : T' \in \Gamma(u)\}$$

for all nodes  $u$  of  $T$ , starting from the leaves, and making its way to the root.

This problem was used by Horn to solve a scheduling problem [Hor72]. In Horn's model, we have jobs represented by nodes of a tree, with the constraint that each job  $u$  can only be executed after the jobs on the path from  $u$  to the root of the tree are executed. These constraints are called precedence constraints. Suppose that job  $u$  takes time  $b_u$ , and that we promised the client who asked for job  $u$  to be executed that we would refund them  $a_u$  dollars for each time unit they have to wait before  $u$  is executed. Our goal is to schedule the jobs one after the other so that we minimize the total refunds, i.e. to arrange the jobs in a sequence  $u_1, \dots, u_n$  so that the objective

$$\sum_{i=1}^n a_{u_i} \sum_{j=1}^i b_{u_j}$$

is minimized. Horn showed that this problem can be solved efficiently if we can find the  $r(u)$  values defined above. Check his paper for the details.

In this note we describe Horn's algorithm for computing the  $r(u)$  values. More specifically, we describe an efficient implementation of it, due to Galil [Gal80], using mergeable priority queues.

Before we describe the algorithm, let's make one simple but crucial observation.

**Proposition 1.** *If  $a, b, c, d$  are positive numbers, then*

$$\frac{a+c}{b+d} \geq \frac{a}{b} \iff \frac{c}{d} \geq \frac{a}{b}.$$

*Similarly, if  $a, b, c, d, c', d'$  are positive numbers, then*

$$\frac{a+c}{b+d} \geq \frac{a+c'}{b+d'} \iff \frac{c}{d} \geq \frac{c'}{d'}.$$


*Proof.* Let  $r = \frac{a}{b}$  and  $r' = \frac{c}{d}$ . Then

$$a+c = r(b+d) + (r'-r)d.$$

Since  $d > 0$ , if  $r' - r \geq 0$ , then  $a+c \geq r(b+d)$ , i.e.  $\frac{a+c}{b+d} \geq \frac{a}{b}$ . Conversely, if  $r' - r < 0$ , then  $a+c < r(b+d)$ , i.e.  $\frac{a+c}{b+d} < \frac{a}{b}$ .

The proof for the second claim is analogous.  $\square$

Consider the figure below, where  $T'$  and  $T''$  are subtrees of some larger tree, and, say,  $T'$  is rooted at  $u^*$ . In this picture, the parent of the root node of  $T''$  lies in  $T'$ . By Proposition 1, if  $r(T'') \geq r(T')$ , then  $r(T' \cup T'') \geq r(T')$ , so if  $T'$  maximizes  $r(T')$  over all subtrees in  $\Gamma(u^*)$ , then so does  $T' \cup T''$ . In other words, if  $r(u^*) = r(T')$  and  $r(T'') \geq r(T')$ , then  $r(u^*) = r(T' \cup T'')$ . Conversely, if  $r(T'') < r(T')$ , then  $T' \cup T''$  cannot be optimal, i.e.  $r(T' \cup T'') < r(T') \leq r(u^*)$ .



This motivates an algorithm which keeps merging trees as long as the  $r$ -value increases. To describe it formally, we need one more piece of notation. For a subtree  $T'$ , we let  $B(T')$  be the set of nodes of  $T$  which are not in  $T'$  but whose parent is in  $T'$ . This is sort of a “boundary” around  $T'$ . (See Figure 2).




Figure 2: The blue nodes are the boundary set  $B(T')$  for the subtree  $T'$  shown in red.

We first describe the algorithm on a high level. Then we will discuss implementation. In the algorithm description below  $\tilde{T}_u$  will eventually equal a subtree rooted at  $u$  such that  $r(\tilde{T}_u) = r(u)$ , i.e. an optimal subtree rooted at  $u$ . If there are many such subtrees,  $\tilde{T}_u$  will be the one with the maximum number of nodes. (This makes it unique, although this fact is not immediately obvious.)

1. For each leaf node  $u$ , set  $\tilde{T}_u$  to just be the node  $u$ , and  $r(u)$  to be  $\frac{a_u}{b_u}$ . (There is no other option for the leaves.) Call the leaf nodes “processed”.
2. Take any internal node  $u$  all of whose children have been processed. Let its children be  $v$  and  $w$ . Initialize  $\tilde{T}_u$  to  $u$ ,  $r(u)$  to  $\frac{a_u}{b_u}$ , and  $B(\tilde{T}_u)$  to  $\{v, w\}$ .
3. As long as  $B(\tilde{T}_u)$  is not empty, take the node  $v \in B(\tilde{T}_u)$  with the largest  $r(v)$ . If  $r(v) < r(u)$  or if  $B(\tilde{T}_u)$  is empty, call  $u$  “processed”. If, otherwise,  $r(v) \geq r(u)$ , then reset  $r(u)$  to  $r(\tilde{T}_u \cup \tilde{T}_v)$ ,  $\tilde{T}_u$  to  $\tilde{T}_u \cup \tilde{T}_v$ , and reset  $B(\tilde{T}_u)$  to  $B(\tilde{T}_u \cup \tilde{T}_v)$ .
4. Repeat steps 2. and 3. until all nodes are processed.

For the correctness of this algorithm, we need to argue that when a node  $u$  has been processed, then  $\tilde{T}_u$  is the subtree rooted at  $u$  which satisfies  $r(\tilde{T}_u) = r(u) = \max\{r(T') : T' \in \Gamma(u)\}$  and has maximum number of nodes. We sketch the argument. It is obvious that this holds for the leaves, because for any leaf node  $u$ ,  $\Gamma(u)$  contains only one subtree: the single-node tree  $u$  itself. For the other nodes, we proceed by induction (on the number of steps of the algorithm). When we are processing a node  $u$  we know that all nodes below it have already been processed. Let  $T^*$  be the subtree rooted at  $u$  such that  $r(T^*) = \max\{r(T') : T' \in \Gamma(u)\}$  and  $T^*$  has the maximum number of nodes among such subtrees. Our goal is to show that when  $u$  is already processed,  $\tilde{T}_u = T^*$ .

We claim that at any point in time when  $u$  is still being processed,  $\tilde{T}_u$  is a subtree of  $T^*$ . This is certainly true initially, when  $\tilde{T}_u$  is just  $u$ , because we know that  $T^*$  must be rooted at  $u$ . We proceed by induction on the number of times we add a tree to  $\tilde{T}_u$ . Suppose that at some point  $B(\tilde{T}_u) = \{u_1, \dots, u_k\}$ , where  $r(u_1) \geq r(u_2) \geq \dots \geq r(u_k)$ . The algorithm adds  $\tilde{T}_{u_1}$  to  $\tilde{T}_u$  if  $r(u_1) \geq r(u)$ . If  $r(u_1) < r(u)$  then there is nothing to prove because  $\tilde{T}_u$  does not change. So, suppose that  $r(u_1) \geq r(u)$  but, towards contradiction,  $\tilde{T}_{u_1}$  is not a subtree of  $T^*$ . Then we have the following claims:

- (a) We must have  $u_1 \notin T^*$ .
- (b)  $r(u_1) = r(\tilde{T}_{u_1}) < r(T^*)$

Both claims hold because, otherwise, we can merge  $\tilde{T}_{u_1}$  with  $T^*$  and, by the second case of Proposition 1, not decrease  $r(T^*)$ . These two claims lead to a contradiction: the first one implies that  $r(T^*) \leq \max\{r(u), r(u_2)\} \leq r(u_1)$  (try to verify this using Proposition 1), which contradicts the second one.

We have then established that  $\tilde{T}_u$  is a subtree of  $T^*$  after  $u$  is processed. It cannot be a strict subtree, because if it were, then there would be some  $v \in B(\tilde{T}_u)$  which is a node of  $T^*$ . But we know that for all  $v$  in  $B(\tilde{T}_u)$  we have  $r(v) < r(\tilde{T}_u)$ , which would imply that  $r(T^*) < r(\tilde{T}_u)$ . (Again, try to verify this using Proposition 1.) Therefore,  $\tilde{T}_u = T^*$ , which completes the analysis.

Let us now see how to implement this algorithm. We will not actually keep the trees  $\tilde{T}_u$  because we only care about the values  $r(u)$ . (You can modify the implementation to also compute the trees without increasing the asymptotic running time: this is a useful exercise.) The idea is that every node  $u$  will have a binomial max-heap  $u.H$  which stores the nodes in  $B(\tilde{T}_u)$ , and the key of every node  $v$  equals  $r(v)$ , which we will store in  $v.r$ . Merging  $\tilde{T}_u$  and  $\tilde{T}_v$  just corresponds to taking the union of  $u.H$  and  $v.H$ . We will also store the sum  $\sum_{v \in \tilde{T}_u} a_v$  in  $u.num$  and the sum  $\sum_{v \in \tilde{T}_u} b_v$  in  $u.denom$ , so that we can easily compute  $r(u) = u.r = \frac{u.num}{u.denom}$ .

The code below takes a pointer  $u$  to a node in  $T$

```

COMPUTER( $u$ )
1   $u.num = a_u$ 
2   $u.denom = b_u$ 
3   $u.r = u.num/u.denom$ 
4  if  $u$  is a leaf
5      Initialize  $u.H$  to be empty and return
6  else Let  $v$  and  $w$  be the children of  $u$ 
7      COMPUTER( $v$ )
8      COMPUTER( $w$ ) // If  $u$  has only one child  $v$ , ignore COMPUTER( $w$ )
9      Initialize  $u.H$  to contain  $v$  and  $w$ , with keys  $v.r$  and  $w.r$ .
10     while  $u.H$  is not empty
11          $v = \text{MAX}(u.H)$ 
12         if  $v.r \geq u.v$ 
13              $u.num = u.num + v.num$ 
14              $u.denom = u.denom + v.denom$ 
15              $u.r = u.num/u.denom$ 
16             EXTRACT-MAX( $u.H$ )
17              $u.H = \text{UNION}(u.H, v.H)$ 
18         else return
19 return

```

To compute all the  $r$ -values, we just call COMPUTER( $u^*$ ) once at the root  $u^*$ . Notice that the procedure is only called once for every root of  $T$ . It is easy to see that the running time is dominated by the total running time of the heap operations. The total number of insertions is bounded by  $n$ , since each node is inserted at most once. The number of MAX and EXTRACT-MAX operations is bounded by the number of UNION operations. We can charge any UNION( $u.H, v.H$ ) operation, where  $v$  is on a lower level than  $u$ , to  $v$ . Every node is charged in this way at most once, because we remove it from the heap  $u.H$  before the union, so  $v$  will never again be returned by a MAX operation and participate in a union. The number of unions is then bounded by  $n$ , and, therefore, the total number of heap operations is bounded by  $O(n)$ . Since each binomial heap operation takes time  $O(\log n)$ , the total running time is  $O(n \log n)$ .

## References

- [Gal80] Zvi Galil. Applications of efficient mergeable heaps for optimization problems on trees. *Acta Inf.*, 13:53–58, 1980.
- [Hor72] W. A. Horn. Single-machine job sequencing with treelike precedence ordering and linear delay penalties. *SIAM Journal on Applied Mathematics*, 23(2):189–202, 1972.