

# Embedded SQL

CSC343 Introduction to Databases

# Problems with using interactive SQL

- Standard SQL is not “Turing-complete”.
  - E.g., Two profs are “colleagues” if they’ve co-taught a course or share a colleague.
  - We can’t write a query to find all colleagues of a given professor because we have no loops or recursion.
- Most DB users shouldn’t be writing SQL queries!
  - We want to run queries that are **based on** user input, not have users writing actual queries.

# SQL + a conventional language

- If we can combine SQL with code in a conventional language, we can solve these problems.
- Issue: SQL is based on relations, and conventional languages have no such type.
- Solution:
  - feed tuples from SQL to the other language one at a time;
  - feed each attribute value into a particular variable.

# Approaches

- Three approaches for combining SQL and a general-purpose language:
  - Stored Procedures
  - Statement-level Interface
  - Call-level interface

# I. Stored Procedures

- The SQL standard includes a language for defining “stored procedures”, which can
  - have parameters and a return value,
  - use local variables, ifs, loops, etc.,
  - execute SQL queries.
- Stored procedures can be used in these ways:
  - called from the interpreter,
  - called from SQL queries,
  - called from another stored procedure,
  - be the action that a trigger performs.

# Example (just to give an idea)

- A boolean function `Comedy(y INT, s CHAR(15))` that returns true iff
  - movie studio `s` produced no movies in year `y`, or
  - produced at least one comedy.
- Reference: Ullman and Widom textbook, chapter 9

## Reference: textbook figure 9.1.3

```
CREATE FUNCTION Comedy(y INT, s CHAR(15)) RETURNS BOOLEAN
IF NOT EXISTS
    (SELECT *
      FROM Movies
      WHERE year = y AND studioName = s)
THEN RETURN TRUE;
ELSIF 1 <=
    (SELECT COUNT(*)
      FROM Movies
      WHERE year = y AND studioName = s AND
            genre = 'comedy')
THEN RETURN TRUE;
ELSE RETURN FALSE;
END IF;
```

# Calling a stored procedure

- ```
SELECT StudioName  
FROM Studios  
WHERE Comedy(2010, StudioName);
```



# Not standard!

- The language is called **SQL/PSM** (Persistent Stored Modules).
  - It came into the SQL standard in SQL3, 1999.
- By then, commercial DBMSs had defined their own proprietary languages for stored procedures
  - They have generally stuck to them.
- PostgreSQL has defined **PL/pgSQL**.
  - It supports some, but not all, of SQL/PSM.
  - Reference: Chapter 39 of the PostgreSQL documentation.

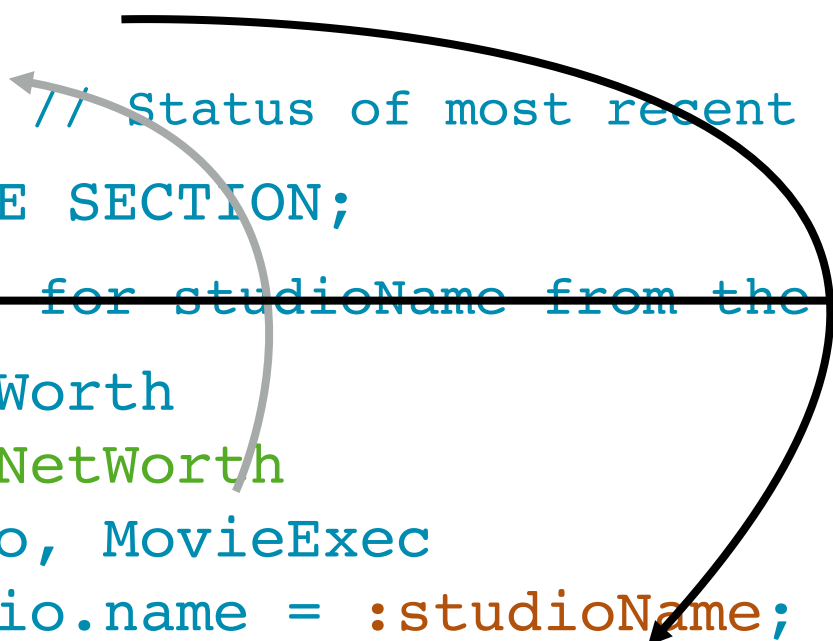
## 2. Statement-level interface (SLI)

- Embed SQL statements into code in a conventional language like C or Java.
- Use a preprocessor to replace the SQL with calls written in the host language to functions defined in an SQL library.
- Special syntax indicates which bits of code the preprocessor needs to convert.

# Example, in C (just to give you an idea)

Reference: textbook example 9.7

```
void printNetWorth() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char studioName[50];  
        int presNetWorth;  
        char SQLSTATE[6]; // Status of most recent SQL stmt  
    EXEC SQL END DECLARE SECTION;  
    /* OMITTED: Get value for studioName from the user. */  
    EXEC SQL SELECT netWorth  
        INTO :presNetWorth  
        FROM Studio, MovieExec  
        WHERE Studio.name = :studioName;  
    /* OMITTED: Report back to the user */  
}
```



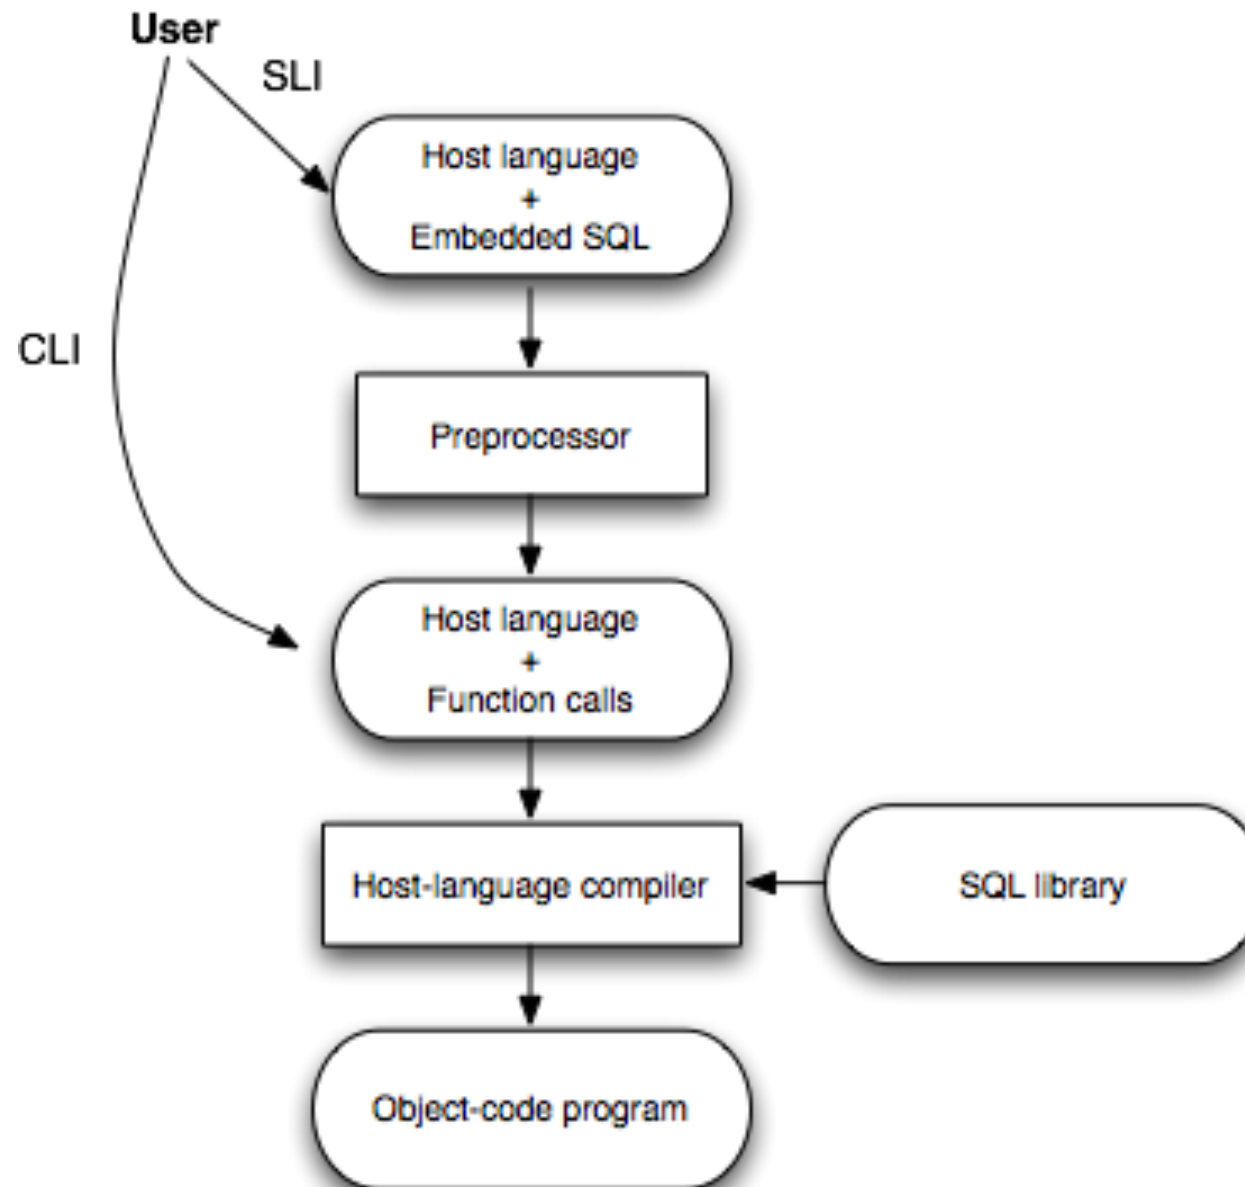
# Data from host language $\Rightarrow$ SQL

- Some special syntax tells the preprocessor things like this:
  - “Variable `studioName` my C code may be referred to by the SQL that’s embedded in my C code.”
- Uses the keywords `DECLARE` `SECTION`.

# Data from SQL $\Rightarrow$ host language

- Again, shared variables make the connection.
- If the query will yield a single tuple, an `INTO` clause can be added to the query.
  - It lists the host-language variables that should receive the attribute values from that one tuple.
- For multi-tuple results, a “cursor” allows you to iterate over the tuples.
  - Much like an iterable object in an Object-oriented language.
  - We won’t cover this, but see figures 9.8 and 9.9 for examples.

# Big picture



### 3. Call-level interface (CLI)

- Instead of using a pre-processor to replace embedded SQL with calls to library functions, we write those calls.
- Eliminates need to preprocess.
- Each language has its own set of library functions for this.
  - for C, it's called SQL/CLI
  - for Java, it's called JDBC
  - for PHP, it's called PEAR DB
- We'll look at just one: JDBC.

JDBC



# JDBC Example (textbook section 9.6)

Do this once in a program:

```
/* Get ready to execute queries. */  
import java.sql.*;  
  
/* A static method of the Class class. It loads the  
   specified driver */  
Class.forName("org.postgresql.jdbc.Driver");  
  
Connection conn = DriverManager.getConnection(  
    jdbc:postgresql://localhost:5432/csc343h-<my user_name>,  
    <user_name>,  
    "" );  
  
/* Continued ... */
```

# The arguments to getConnection

- `jdbc:postgresql`

We'll use this, but it could be, e.g., `jdbc:mysql`

- `localhost:5432`

Use this for the CS Teaching Labs.

- `csc343h-userid` and `userid`

Substitute your `userid` on the CS Teaching Labs.

- `""`

Password (unrelated to your password).

Literally use the empty string.

# Do this once per query in a program:

```
/*Execute a query and iterate through the resulting tuples.*/
```

```
PreparedStatement execStat = conn.prepareStatement(  
    "SELECT netWorth FROM MovieExec");
```

```
ResultSet worths = execStat.executeQuery();
```

```
while (worths.next()) {
```

```
    int worth = worths.getInt(1);
```

```
    /* If the tuple also had a float and another int  
       attribute, you'd get them by calling  
       worths.getFloat(2) and worths.getInt(3).  
       Or you can look up values by attribute name.  
       Example: worths.getInt(netWorth)
```

```
    */
```

```
    /* OMITTED: Process this net worth */
```

```
}
```

# The Java details

- For full details on the Java classes and methods used, see the Java API documentation:

<https://docs.oracle.com/javase/8/docs/api/java/sql/package-summary.html>

# Exceptions can occur

- Any of these calls can generate an exception, therefore, they should be inside try/catch blocks.

```
try {  
    /* OMITTED: JDBC code */  
}  
  
catch (SQLException ex) {  
    /* OMITTED: Handle the exception */  
}
```

- The class `SQLException` has methods to return the `SQLSTATE`, etc.

# Prepare separately vs execute immediately

- We can combine preparation and execution.

- Separate looked like this:

```
PreparedStatement pStat =  
    conn.prepareStatement(  
        "SELECT netWorth FROM MovieExec");  
ResultSet worths =  
    pStat.executeQuery();
```

- Combined looks like this:

```
Statement stat =  
    conn.createStatement();  
ResultSet worths =  
    stat.executeQuery()
```

# What is “preparation”?

- Preparing a statement includes:
  - parsing the SQL
  - compiling
  - optimizing
- The resulting `PreparedStatement` can be executed any number of times without having to repeat these steps.

# If the query isn't known until run time

- We may need input and computation to determine exactly what the query should be.
- In that case:
  - Hard-code in the parts known.
  - Use the character `?` as a placeholder for the values not known. (Don't put it in quotes!)
- This is enough to allow a `PreparedStatement` to be constructed.
- Once all values for the placeholders are known, use methods `setString`, `setInt`, etc., to fill in those values.



# Example (Figure 9.22)

```
PreparedStatement studioStat =  
    conn.prepareStatement(  
        "INSERT INTO Studio(name, address)  
        VALUES(?, ?)"  
    );  
  
/* OMITTED: Get values for studioName and studioAddr */  
studioStat.setString(1, studioName);  
  
studioStat.setString(2, studioAddr);  
  
studioStat.executeUpdate();
```

# Why not just build the query in a string?

- We constructed an incomplete `PreparedStatement` and filled in the missing values using method calls.
- Instead, we could just build up the query in an ordinary string at run time, and ask to execute that.
- There are classes and methods that will do this in JDBC.

# Example that builds the query in a string

- We can just use a `Statement`, and give it a `String` to execute.

```
// stat cannot be compiled & optimized (yet).  
Statement stat = conn.createStatement();
```

```
String query =  
    "SELECT networth  
    FROM MovieExec  
    WHERE execName like '%Spielberg%';  
    "
```

```
// executeQuery can now compile, optimize and run  
// the query.  
ResultSet worths = stat.executeQuery(query);
```

# What could go wrong?

# Example: vulnerable code

Suppose we want the user to provide the string to compare to.

```
Statement stat = conn.createStatement();
```

```
String who = /* get a string from the user */
```

```
String query =  
    "SELECT networth  
    FROM MovieExec  
    WHERE execName like '%" + who + "%';  
    "
```

```
ResultSet worths = stat.executeQuery(query);
```

If a user enters **Milch**, the SQL code executed is this:

```
SELECT networth  
FROM MovieExec  
WHERE execName like '%Milch%';
```

Nothing bad happens.

# An injection can exploit the vulnerability

What could a malicious user enter?

```
SELECT networth  
FROM MovieExec  
WHERE execName like '%?????????????%';
```

# An injection can exploit the vulnerability

But if a malicious user enters

```
Milch%'; drop table Contracts; --
```

the code we execute is this:

```
SELECT networth  
FROM MovieExec  
WHERE execName like '%Milch%'; DROP TABLE Contracts; --%';
```

In other words:

```
SELECT networth  
FROM MovieExec  
WHERE execName like '%Milch%';
```

```
DROP TABLE Contracts; --%';
```



HI, THIS IS  
YOUR SON'S SCHOOL.  
WE'RE HAVING SOME  
COMPUTER TROUBLE.



OH, DEAR - DID HE  
BREAK SOMETHING?  
IN A WAY-



DID YOU REALLY  
NAME YOUR SON  
Robert'); DROP  
TABLE Students;-- ?



OH, YES. LITTLE  
BOBBY TABLES,  
WE CALL HIM.

WELL, WE'VE LOST THIS  
YEAR'S STUDENT RECORDS.  
I HOPE YOU'RE HAPPY.



AND I HOPE  
YOU'VE LEARNED  
TO SANITIZE YOUR  
DATABASE INPUTS.

Reference: <https://xkcd.com/327/>

# Always use a PreparedStatement

- This was an example of an **injection**.
- The simple approach of giving a String to a Statement is vulnerable to injections.
- Moral of the story: always use a PreparedStatement.

# Queries vs updates in JDBC

- The previous examples used `executeQuery`.
- This method is only for pure queries.
- For SQL statements that change the database (insert, delete or modify tuples, or change the schema), use the analogous method `executeUpdate`.