

SQL:

Data Manipulation Language

CSC343, Introduction to Databases



UNIVERSITY OF
TORONTO

Introduction

- So far, we have defined database schemas and queries mathematically.
- SQL is a formal language for doing so with a DBMS.
- “Structured Query Language”, but it’s for more than writing queries.
- Two sub-parts:
 - DDL (Data Definition Language), for defining schemas.
 - DML (Data Manipulation Language), for writing queries and modifying the database.

PostgreSQL

- We'll be working in PostgreSQL, an open-source relational DBMS.
- SQL Standards?
 - There are several, the most recent being SQL:2008.
 - The standards are not freely available. Must purchase from the International Standards Organization (ISO).
 - PostgreSQL supports most of it SQL:2008.
 - DBMSs vary in some details, making portability difficult.

SQL - A high-level language

- SQL is a very high-level language.
 - Say “what” rather than “how.”
- We write queries without manipulating data – in contrast with languages like Java or C++.
- SQL provides physical “data independence:”
 - Details of how the data is stored can change with no impact on queries.
- You can focus on readability, but the DMBS optimizes our query, so we get efficiency.

SELECT vs σ

- In SQL,
 - “SELECT” is for choosing columns, *i.e.*, Π .
 - Example:

```
SELECT surName  
FROM Student  
WHERE campus = 'StG';
```

- In relational algebra,
 - “select” means choosing rows, *i.e.*, σ .

Meaning of a query with one relation

```
SELECT name  
FROM Course  
WHERE dept = 'CSC';
```

$$\pi_{\text{name}} (\sigma_{\text{dept}=\text{"csc"}} (\text{Course}))$$

Meaning of a query with one relation

```
SELECT name  
FROM Course  
WHERE dept = 'CSC';
```



(Course)

($\sigma_{\text{dept}=\text{'csc'}}$ (Course))

π_{name} ($\sigma_{\text{dept}=\text{'csc'}}$ (Course))

... and with multiple relations

```
SELECT name  
FROM Offering, Took  
WHERE Offering.id = Took.oid and  
      dept = 'CSC';
```

$$\pi_{\text{name}} (\sigma_{\text{Offering.id=Took.id} \wedge \text{dept='csc'}} (\text{Offering} \times \text{Took}))$$

Temporarily renaming a table

- You can rename tables (just for the duration of the statement):

```
SELECT e.name, d.name  
FROM employee e, department d  
WHERE d.name = 'marketing'  
AND e.name = 'Horton';
```

- Can be convenient vs the longer full names:

```
SELECT employee.name, department.name  
FROM employee, department  
WHERE department.name = 'marketing'  
AND employee.name = 'Horton';
```

Self-joins

- renaming is *required* for self-joins.

- Example:

-

```
SELECT e1.name, e2.name
FROM   employee e1, employee e2
WHERE  e1.salary < e2.salary;
```

* In SELECT clauses

- A * in the SELECT clause means “all attributes of this relation.”

- Example:



```
SELECT *  
FROM   Course  
WHERE  dept = 'CSC';
```

Renaming attributes

- Use `AS «new name»` to rename an attribute in the result.

- Example:

- ```
SELECT name AS title, dept
FROM Course
WHERE breadth;
```

# Complex Conditions in a WHERE

- We can build boolean expressions with operators that produce boolean results.
  - comparison operators: `=`, `<>`, `<`, `>`, `<=`, `>=`
  - and many other operators:  
see section 6.1.2 of the text and chapter 9 of the postgresSQL documentation.
- Note that “not equals” is: `<>`
- We can combine boolean expressions with:  
Boolean operators: `AND`, `OR`, `NOT`.

# Example: Compound condition

- Find 3rd- and 4th-year CSC courses:

```
SELECT *
FROM Offering
WHERE dept = 'CSC' AND cnum >= 300;
```

# ORDER BY

- To put the tuples in order, add this as the final clause:  
`ORDER BY «attribute list» [DESC]`
- 
- The default is ascending order; DESC overrides it to force descending order.
- The attribute list can include expressions: e.g., `ORDER BY sales+rentals`
- The ordering is the last thing done before the SELECT, so all attributes are still available.

# Case-sensitivity and whitespace

- Example query:

```
SELECT surName
FROM Student
WHERE campus = 'StG' ;
```

- Keywords, like `SELECT`, are not case-sensitive.
- Identifiers, like `Student` are not case-sensitive either.
- Literal strings, like `'StG'`, are case-sensitive, and require single quotes.
- Whitespace (other than inside quotes) is ignored.



# Expressions in SELECT clauses

- Instead of a simple attribute name, you can use an expression in a SELECT clause.
- Operands: attributes, constants  
Operators: arithmetic ops, string ops
- Examples:  

```
SELECT sid, grade+10 AS adjusted
FROM Took;
```

```
SELECT dept || cnum
FROM course;
```

# Expressions that are a constant

- Sometimes it makes sense for the whole expression to be a constant (something that doesn't involve any attributes!).
- Example:  

```
SELECT dept, cNum,
 'satisfies' AS breadthRequirement
FROM Course
WHERE breadth;
```

# Pattern operators

- Two ways to compare a string to a pattern by:
  - «*attribute*» LIKE «*pattern*»
  - «*attribute*» NOT LIKE «*pattern*»
- Pattern is a quoted string
  - % means: any string
  - \_ means: any single character
- Example:

```
SELECT *
FROM Course
WHERE name LIKE '%Comp%';
```

# Aggregation

# Computing on a column

- We often want to compute something across the values in a column.
- `SUM`, `AVG`, `COUNT`, `MIN`, and `MAX` can be applied to a column in a `SELECT` clause.
- Also, `COUNT ( * )` counts the number of tuples.
- We call this aggregation.
- Note: To stop duplicates from contributing to the aggregation, use `DISTINCT` inside the brackets. (Does not affect `MIN` or `MAX`.)

# Grouping

- If we follow a SELECT-FROM-WHERE expression with GROUP BY <attributes>
  - The tuples are grouped according to the values of those attributes, and
  - any aggregation gives us a single value per group.

# Restrictions on aggregation

- If any aggregation is used, then each element of the SELECT list must be either:
  - aggregated, or
  - an attribute on the GROUP BY list.
- Otherwise, it doesn't even make sense to include the attribute.

# HAVING Clauses

- **Example:** having.txt
- WHERE let's you decide which tuples to keep.
- Similarly, you can decide which *groups* to keep.
- Syntax:
  - . . .
  - GROUP BY «*attributes*»
  - HAVING «*condition*»
- Semantics:
  - Only groups satisfying the condition are kept.




# Restrictions on HAVING clauses

- Outside subqueries, HAVING may refer to attributes only if they are either:
  - aggregated, or
  - an attribute on the GROUP BY list.
- (Same requirement as for SELECT clauses with aggregation)

# Order of execution of a SQL query

| Query order | Execution order |
|-------------|-----------------|
| SELECT      | FROM            |
| FROM        | WHERE           |
| WHERE       | GROUP BY        |
| GROUP BY    | HAVING          |
| HAVING      | SELECT          |
| ORDER BY    | ORDER BY        |



# Set operations

# Tables can have duplicates in SQL

- A table can have duplicate tuples, unless this would violate an integrity constraint.
- And SELECT-FROM-WHERE statements leave duplicates in unless we say not to.
- Why?
  - Getting rid of duplicates is expensive!
  - We may want the duplicates because they tell us how many times something occurred.

# Relational Algebra with Bags

- Behaviour of most operations is no different.
- - $\sigma$ ,  $\rho$ : as before
  - $\pi$ : duplicates are not removed.
  - joins: duplicates can proliferate

# Bags

- SQL treats tables as “bags” (or “multisets”) rather than sets.
- Bags are just like sets, but duplicates are allowed.
- $\{6, 2, 7, 1, 9\}$  is a set (and a bag)  
 $\{6, 2, 2, 7, 1, 9\}$  is not a set, but is a bag.
- Like with sets, order doesn't matter.  
 $\{6, 2, 7, 1, 9\} = \{1, 2, 6, 7, 9\}$

# Union, Intersection, and Difference

- These are expressed as:

( «*subquery*» ) UNION ( «*subquery*» )

( «*subquery*» ) INTERSECT ( «*subquery*» )

( «*subquery*» ) EXCEPT ( «*subquery*» )

- The brackets are mandatory.
- The operands must be queries; you can't simply use a relation name.

# Example

```
(SELECT sid
 FROM Took
 WHERE grade > 95)
 UNION
(SELECT sid
 FROM Took
 WHERE grade < 50);
```



# Operations $\cup$ , $\cap$ , and $-$ with Bags

- For  $\cup$ ,  $\cap$ , and  $-$  the number of occurrences of a tuple in the result requires some thought.

$$1. \{1, 1, 1, 3, 7, 7, 8\} \cup \{1, 5, 7, 7, 8, 8\}$$

$$= \{1, 1, 1, 3, 7, 7, 8, 1, 5, 7, 7, 8, 8\}$$

$$= \{1, 1, 1, 1, 3, 5, 7, 7, 7, 7, 8, 8, 8\}$$

$$2. \{1, 1, 1, 3, 7, 7, 8\} \cap \{1, 5, 7, 7, 8, 8\}$$

$$= \{1, 7, 7, 8\}$$

$$3. \{1, 1, 1, 3, 7, 7, 8\} - \{1, 5, 7, 7, 8, 8\}$$

$$= \{1, 1, 3\}$$

# Operations $\cup$ , $\cap$ , and $-$ with Bags

- Suppose tuple  $t$  occurs
  - $m$  times in relation  $R$ , and
  - $n$  times in relation  $S$ .

| Operation  | Number of occurrences of $t$ in result |
|------------|----------------------------------------|
| $R \cap S$ | $\min(m, n)$                           |
| $R \cup S$ | $m + n$                                |
| $R - S$    | $\max(m-n, 0)$                         |

# Bag vs Set Semantics: which is used

- We saw that a **SELECT-FROM-WHERE** statement uses bag semantics by default.
  - Duplicates are kept in the result.
- The set operations use set semantics by default.
  - Duplicates are *eliminated* from the result.

# Motivation: Efficiency

- When doing projection, it is easier not to eliminate duplicates.
  - Just work one tuple at a time.
- For intersection or difference, it is most efficient to sort the relations first.
  - At that point you may as well eliminate the duplicates anyway.

# Controlling Duplicate Elimination

- We can force the result of a SFV query to be a set by using `SELECT DISTINCT ...`
- We can force the result of a set operation to be a bag by using `ALL`, e.g.,

```
(SELECT sid
 FROM Took
 WHERE grade > 95)
 UNION ALL
(SELECT sid
 FROM Took
 WHERE grade < 50);
```

- **Examples:** controlling-dups.txt, except-all.txt

# Views

# The idea

- A view is a relation defined in terms of stored tables (called base tables) and other views.
- Two kinds of views:
  - **Virtual**: no tuples are stored; view is just a query for constructing the relation when needed.
  - **Materialized**: actually constructed and stored.  
Expensive to maintain!
- We'll use only virtual views.



# Example: defining a virtual view

- A view for students who earned an 80 or higher in a CSC course.

```
CREATE VIEW topresults AS
SELECT firstname, surname, cnum
FROM Student, Took, Offering
WHERE
 Student.sid = Took.sid AND
 Took.oid = Offering.oid AND
 grade >= 80 AND dept = 'CSC';
```

# Uses for views

- Break down a large query.
- Provide another way of looking at the same data, e.g., for one category of user.

# Outer Joins

# The joins you know from RA

These can go in a FROM clause:

| Expression                         | Meaning                   |
|------------------------------------|---------------------------|
| <code>R, S</code>                  | $R \times S$              |
| <code>R cross join S</code>        |                           |
| <code>R natural join S</code>      | $R \bowtie S$             |
| <code>R join S on Condition</code> | $R \bowtie_{condition} S$ |

# In practice, natural join is brittle

- A working query can be broken by adding a column to a schema.
  - Example:

```
SELECT sID, instructor
FROM Student NATURAL JOIN Took
 NATURAL JOIN Offering;
```
  - What if we add a column called `campus` to `Offering`?
- Also, having implicit comparisons impairs readability.

Students(sID, surName, campus)

Courses(cID, cName, WR)

Offerings(oID, cID, term, instructor, campus)

Took(sID, oID, grade)

```
SELECT sID, instructor
FROM Student NATURAL JOIN Took
 NATURAL JOIN Offering;
```

# Dangling tuples

- With joins that require some attributes to match, tuples lacking a match are left out of the results. We say that they are “dangling”.
- An **outer join** preserves dangling tuples by padding them with **NULL** in the other relation.
- A join that doesn't pad with **NULL** is called an **inner join**.

# Three kinds of outer join

- **LEFT OUTER JOIN**
  - Preserves dangling tuples from the relation on the LHS by padding with nulls on the RHS.
- **RIGHT OUTER JOIN**
  - The reverse.
- **FULL OUTER JOIN**
  - Does both.



# Example: joining R and S various ways

R

| A | B |
|---|---|
| 1 | 2 |
| 4 | 5 |

S

| B | C |
|---|---|
| 2 | 3 |
| 6 | 7 |

R NATURAL JOIN S

| A | B | C |
|---|---|---|
| 1 | 2 | 3 |

# Example

R

| A | B |
|---|---|
| 1 | 2 |
| 4 | 5 |

S

| B | C |
|---|---|
| 2 | 3 |
| 6 | 7 |

R NATURAL FULL JOIN S

| A    | B | C    |
|------|---|------|
| 1    | 2 | 3    |
| 4    | 5 | NULL |
| NULL | 6 | 7    |

# Example

R

| A | B |
|---|---|
| 1 | 2 |
| 4 | 5 |

S

| B | C |
|---|---|
| 2 | 3 |
| 6 | 7 |

R NATURAL **LEFT** JOIN S

| A | B | C    |
|---|---|------|
| 1 | 2 | 3    |
| 4 | 5 | NULL |

# Example

R

| A | B |
|---|---|
| 1 | 2 |
| 4 | 5 |

S

| B | C |
|---|---|
| 2 | 3 |
| 6 | 7 |

R NATURAL RIGHT JOIN S

| A    | B | C |
|------|---|---|
| 1    | 2 | 3 |
| NULL | 6 | 7 |

# Summary of join expressions

## Cartesian product

`A CROSS JOIN B`

same as `A, B`

## Theta-join

`A JOIN B ON C`

✓ `A {LEFT|RIGHT|FULL} JOIN B ON C`

## Natural join

`A NATURAL JOIN B`

✓ `A NATURAL {LEFT|RIGHT|FULL} JOIN B`

✓ indicates that tuples are padded when needed.

# Keywords INNER and OUTER

- There are keywords **INNER** and **OUTER**, but you never need to use them.
- Your intentions are clear anyway:
  - You get an outer join iff you use the keywords **LEFT**, **RIGHT**, or **FULL**.
  - If you don't use the keywords **LEFT**, **RIGHT**, or **FULL** you get an inner join.

# Impact of having null values

# Missing Information

- Two common scenarios:
  - Missing value.  
E.g., we know a student has some email address, but we don't know what it is.
  - Inapplicable attribute.  
E.g., the value of attribute spouse is inapplicable for an unmarried person.



# Representing missing information

- One possibility: use a special value as a placeholder. E.g.,
  - If age unknown, use 0.
  - If StNum unknown, use 999999999.
- Better solution: use a value not in any domain. We call this a null value.
- Tuples in SQL relations can have **NULL** as a value for one or more components.

# Checking for null values

- You can compare an attribute value to **NULL** with
  - **IS NULL**
  - **IS NOT NULL**
- **Example:**

```
SELECT *
FROM Course
WHERE breadth IS NULL;
```

# In SQL we have 3 truth-values

- Because of **NULL**, we need three truth-values:
  - If one or both operands to a comparison is **NULL**, the comparison *always* evaluates to **UNKNOWN**.
  - Otherwise, comparisons evaluate to **TRUE** or **FALSE**.

# Combining truth values

- We need to know how the three truth-values combine with **AND**, **OR** and **NOT**.
- Can think of it in terms of the truth table.
- Or can think in terms of numbers:
  - **TRUE** = 1, **FALSE** = 0, **UNKNOWN** = 0.5
  - **AND** is min, **OR** is max,
  - **NOT** x is  $(1-x)$ , i.e., it “flips” the value

# The three-valued truth table

| A        | B | A and B | A or B |
|----------|---|---------|--------|
| T        | T | T       | T      |
| TF or FT |   | F       | T      |
| F        | F | F       | F      |
| TU or UT |   | U       | T      |
| FU or UF |   | F       | U      |
| U        | U | U       | U      |

| A | not A |
|---|-------|
| T | F     |
| F | T     |
| U | U     |

# Thinking of the truth-values as numbers

| A        | B | as nums  | A <b>and</b> B | min | A <b>or</b> B | max |
|----------|---|----------|----------------|-----|---------------|-----|
| T        | T | 1, 1     | T              | 1   | T             | 1   |
| TF or FT |   | 1, 0     | F              | 0   | T             | 1   |
| F        | F | 0, 0     | F              | 0   | F             | 0   |
| TU or UT |   | 1, 0.5   | U              | 0.5 | T             | 1   |
| FU or UF |   | 0, 0.5   | F              | 0   | U             | 0.5 |
| U        | U | 0.5, 0.5 | U              | 0.5 | U             | 0.5 |

# Thinking of the truth-values as numbers

| A | as a num, x | not A | 1 - x |
|---|-------------|-------|-------|
| T | 1           | F     | 0     |
| F | 0           | T     | 1     |
| U | 0.5         | U     | 0.5   |

# Surprises from 3-valued logic

- Some laws you are used to still hold in three-valued logic. For example,
  - **AND** is commutative.
- But others don't. For example,
  - The law of the excluded middle breaks:  
 $(p \text{ or } (\text{NOT } p))$  might not be **TRUE**!
  - $(0 * x)$  might not be 0.



# Impact of null values on WHERE

- A tuple is in a query result **iff the WHERE clause is TRUE.**
- UNKNOWN is not good enough.
- “WHERE is picky.”

# Aggregation ignores nulls

|                       | some nulls in A  | All nulls in A |
|-----------------------|------------------|----------------|
| <code>min(A)</code>   |                  | null           |
| <code>max(A)</code>   |                  |                |
| <code>sum(A)</code>   |                  |                |
| <code>avg(A)</code>   |                  |                |
| <code>count(A)</code> |                  | 0              |
| <code>count(*)</code> | all tuples count |                |

# More re the impact of null values

- Other corner cases to think about:
  - `SELECT DISTINCT`: are 2 `NULL` values equal?
  - natural join: are 2 `NULL` values equal?
  - set operations: are 2 `NULL` values equal?
- And later, when we learn about constraints:
  - `UNIQUE` constraint: do 2 `NULL` values violate?
- This behaviour may vary across DBMSs.

# Summary re: NULL

- Any comparison with **NULL** yields **UNKNOWN**.
- **WHERE** is picky: it only accepts **TRUE**.
- Therefore **NATURAL JOIN** is picky too.
- Aggregation ignores **NULL**.
- In other situations where **NULLs** matter
  - when a truth-value may be **NULL**
  - when it matters whether two **NULL** are considered the same

Don't assume. Behaviour may vary by DBMS.

# Subqueries

# Where can a subquery go?

- Relational algebra syntax is so elegant that it's easy to see where subqueries can go.
- In SQL, a bit more thought is required . . .

# Subqueries in a FROM clause

- In place of a relation name in the FROM clause, we can use a subquery.
- The subquery must be parenthesized.
- Must name the result, so you can refer to it in the outer query.

# Worksheet, Q1:

```
SELECT sid, dept || cnum as course, grade
FROM Took,
 (SELECT *
 FROM Offering
 WHERE instructor='Horton') Hofferings
WHERE Took.oid = Hofferings.oid;
```

- This FROM is analogous to:  
 $\text{Took} \times \rho_{\text{Hofferings}} (\text{«subquery»})$



# Subquery as a value in a WHERE

- If a subquery is guaranteed to produce exactly one tuple, then the subquery can be used as a value.
- Simplest situation: that one tuple has only one component.

## Worksheet, Q2:

```
SELECT sid, surname
FROM Student
WHERE cgpa >
 (SELECT cgpa
 FROM Student
 WHERE sid = 99999);
```

- We can't do the analogous thing in RA:

$\Pi_{\text{sid, surname}} \sigma_{\text{cgpa} > (\text{«subquery»})} \text{Student}$

# Special cases

- What if the subquery returns **NULL**?
- What if the subquery could return more than one value?

# Quantifying over multiple results

- When a subquery can return multiple values, we can make comparisons using a quantifier.

- Example:

```
SELECT sid, surname
FROM Student
WHERE cgpa >
 (SELECT cgpa
 FROM Student
 WHERE campus = 'StG');
```

- We can require that
  - cgpa > **all** of them, or
  - cgpa > **at least one** of them.

# The Operator ANY

- Syntax:

$x \text{ «comparison» ANY («subquery»)$

or equivalently

$x \text{ «comparison» SOME («subquery»)$

- Semantics:

Its value is true iff the comparison holds for at least one tuple in the subquery result, i.e.,

$\exists y \in \text{«subquery results»} \mid x \text{ «comparison» } y$

- $x$  can be a *list* of attributes,  
but this feature is not supported by psql.

# The Operator ALL

- Syntax:

$x \text{ «comparison» ALL («subquery»)$

- Semantics:

Its value is true iff the comparison holds for every tuple in the subquery result, i.e.,

$\forall y \in \text{«subquery results»} \mid x \text{ «comparison» } y$

- $x$  can be a list of attributes, but this feature is not supported by psql.

## Universal quantifier:

$x$  «*comparison*» ALL («*subquery*»)

True iff the comparison holds for every row in the subquery result, i.e.,

$$\forall y \in \text{«subquery results»} \mid x \text{ «comparison» } y$$

## Existential quantifier:

$x$  «*comparison*» SOME («*subquery*»)

True iff the comparison holds for at least one row in the subquery result, i.e.,

$$\exists y \in \text{«subquery results»} \mid x \text{ «comparison» } y$$

**x «comparison» ALL («subquery»)**

$\forall y \in \text{«subquery results»} \mid x \text{ «comparison» } y$

**x «comparison» SOME («subquery»)**

$\exists y \in \text{«subquery results»} \mid x \text{ «comparison» } y$

**x IN («subquery»)**

Same as **x = SOME («subquery»)**

**x NOT IN («subquery»)**

Same as **x <> ALL («subquery»)**

} just for  
convenience

**EXISTS («subquery»)**

  $\exists y \in \text{«subquery results»}$



# The Operator IN

- Syntax:

$x$  IN («*subquery*»)

- Semantics:

Its value is true iff  $x$  is in the set of rows generated by the subquery.

- $x$  can be a list of attributes, and psql does support this feature.

## Worksheet, Q3:

```
SELECT sid, dept || cnum AS course, grade
FROM Took NATURAL JOIN Offering
WHERE
 grade >= 80 AND
 (cnum, dept) IN (
 SELECT cnum, dept
 FROM Took NATURAL JOIN Offering
 NATURAL JOIN Student
 WHERE surname = 'Lakemeyer');
```

## Worksheet, Q4:

Suppose we have tables  $R(a, b)$  and  $S(b, c)$ .

I. What does this query do?

```
SELECT a
FROM R
WHERE b IN (SELECT b FROM S);
```

# The Operator EXISTS

- Syntax:  
EXISTS («*subquery*»)
- Semantics:  
Its value is true iff the subquery has at least one tuple.
- Read it as “exists a row in the subquery result”

# Example: EXISTS

```
SELECT surname, cgpa
FROM Student
WHERE EXISTS (
 SELECT *
 FROM Took
 WHERE Student.sid = Took.sid and
 grade > 85);
```

## Worksheet, Q5:

```
SELECT instructor
FROM Offering Off1
WHERE NOT EXISTS (
 SELECT *
 FROM Offering
 WHERE
 oid <> Off1.oid AND
 instructor = Off1.instructor);
```

## Worksheet, Q6:

```
SELECT DISTINCT oid
FROM Took
WHERE EXISTS (
 SELECT *
 FROM Took t, Offering o
 WHERE
 t.oid = o.oid AND
 t.oid <> Took.oid AND
 o.dept = 'CSC' AND
 took.sid = t.sid);
```

# Scope

- Queries are evaluated from the inside out.
- If a name might refer to more than one thing, use the most closely nested one.
- If a subquery refers only to names defined inside it, it can be evaluated **once** and used repeatedly in the outer query.
- If it refers to any name defined outside of itself, it must be evaluated **once for each tuple in the outer query**.

These are called **correlated subqueries**.



# Renaming can make scope explicit

```
SELECT instructor
FROM Offering Off1
WHERE NOT EXISTS (
 SELECT *
 FROM Offering Off2
 WHERE
 Off2.oid <> Off1.oid AND
 Off2.instructor = Off1.instructor);
```

# Summary: where subqueries can go

- As a relation in a FROM clause.
- As a value in a WHERE clause.
- With ANY, ALL, IN or EXISTS in a WHERE clause.
- As operands to UNION, INTERSECT or EXCEPT.

# Modifying a Database

# Database Modifications

- Queries return a relation.
- A modification command does not; it changes the database in some way.
- Three kinds of modifications:
  - Insert a tuple or tuples.
  - Delete a tuple or tuples.
  - Update the value(s) of an existing tuple or tuples.

# Two ways to insert



INSERT INTO «*table*» VALUES «*list of rows*»;

INSERT INTO «*table*» («*subquery*»);

# Naming attributes in INSERT

- Sometimes we want to insert tuples, but we don't have values for all attributes.
- If we name the attributes we *are* providing values for, the system will use **NULL** or a default for the rest.

# Example

```
CREATE TABLE Invite (
 name TEXT,
 campus TEXT DEFAULT 'StG',
 email TEXT,
 age INT);
```

```
INSERT INTO Invite(name, email)
(SELECT firstname, email
 FROM Student
 WHERE cgpa > 3.4);
```

Here, name and email get values from the query, campus gets the default value, and age gets **NULL**.

# Deletion

- Delete tuples satisfying a condition:

```
DELETE FROM «relation»
WHERE «condition»;
```

- Delete all tuples:

```
DELETE FROM «relation»;
```



# Example 1: Delete Some Tuples

```
DELETE FROM Course
WHERE NOT EXISTS (
 SELECT *
 FROM Took JOIN Offering
 ON Took.oid = Offering.oid
 WHERE
 grade > 50 AND
 Offering.dept = Course.dept AND
 Offering.cnum = Course.cnum
);
```

# Updates

- To change the value of certain attributes in certain tuples to given values:

```
UPDATE «relation»
SET «list of attribute assignments»
WHERE «condition on tuples»;
```

# Example: update one tuple

- Updating one tuple:

```
UPDATE Student
SET campus = 'UTM'
WHERE sid = 99999;
```

- Updating several tuples:

```
UPDATE Took
SET grade = 50
WHERE grade >= 47 and grade < 50;
```

# Updates on Views

- **Generally, it is impossible to modify a virtual view, because it doesn't exist.**
- Can't we “translate” updates on views into “equivalent” updates on base tables?
  - Not always (in fact, not often).
  - Most systems prohibit most view updates.

# Example: The View

- CREATE VIEW Synergy AS
- SELECT Likes.drinker, Likes.beer, Sells.bar
- FROM Likes, Sells, Frequents
- WHERE Likes.drinker = Frequents.drinker
- AND Likes.beer = Sells.beer
- AND Sells.bar = Frequents.bar;

Pick one copy of  
each attribute

Natural join of Likes,  
Sells, and Frequents

# Interpreting a View Insertion

- We cannot insert into Synergy --- it is a virtual view.
- But we could try to translate a (drinker, beer, bar) triple into three insertions of projected pairs, one for each of Likes, Sells, and Frequent.
  - Sells.price will have to be **NULL**.
  - There isn't always a unique translation.

# Materialized Views

- Problem: each time a base table changes, the materialized view may change.
  - Cannot afford to recompute the view with each change.
- Solution: Periodic reconstruction of the materialized view, which is otherwise “out of date.”

# Example: A Data Warehouse

- Wal-Mart stores every sale at every store in a database.
- Overnight, the sales for the day are used to update a data warehouse = materialized views of the sales.
- The warehouse is used by analysts to predict trends and move goods to where they are selling best.