

1. **k -COLOUR $\in NP$:** Consider the following verifier:

On input (G, c) , where c is an assignment of colours to the vertices of G , check that c uses no more than k colours and that every edge of G has endpoints with different colours.

This runs in polytime: a linear-time loop over the edges of G and a linear-time loop over the values of c .

Moreover, the verifier outputs TRUE for some c iff there exists a colouring of G using no more than k colours.

k -COLOUR is NP-hard: For an arbitrary $k \geq 3$, we show $3\text{-COLOUR} \rightarrow_p k\text{-COLOUR}$.

On input G , output G' equal to G with the addition of $k - 3$ new vertices u_1, \dots, u_{k-3} , with edges between every pair of new vertices and between every new vertex and every vertex in G .

Clearly, G' can be computed in polytime from G : we add a constant number of new vertices and a linear number of new edges, with a simple, regular structure.

Also, if G can be coloured using at most 3 colours, then G' can be coloured using at most k colours: use 3 colours for the vertices of G and $k - 3$ other colours for u_1, \dots, u_{k-3} (one colour for each vertex).

Finally, if G' can be coloured using at most k colours, then u_1, \dots, u_{k-3} must be assigned $k - 3$ different colours (because u_1, \dots, u_{k-3} forms a clique), and none of those colours can be assigned to any vertex in G (because each u_i is adjacent to every vertex of G). This means the vertices of G can be coloured using only the 3 remaining colours.

ALTERNATE REDUCTION

Show that $k\text{-COLOUR} \rightarrow_p (k + 1)\text{-COLOUR}$ for every $k \geq 3$, (then $k\text{-COLOUR}$ is NP-hard for every $k \geq 3$, by induction):

On input G , output G' equal to G with the addition of one new vertex v_0 together with edges between v_0 and every vertex in G .

Clearly, G' can be computed in linear time from G .

Also, if G can be coloured with at most k colours, then G' can be coloured using at most $k + 1$ colours: we only need one more colour for v_0 .

Finally, if G' can be coloured using at most $k + 1$ colours, then the colour assigned to v_0 must be different from the colour of every other vertex. This means the vertices of G are coloured using at most k colours.

2. **Search Problem:** The 3-COLOUR-SEARCH problem can be defined as follows.

- **In:** Undirected graph $G = (V, E)$.
- **Out:** Assignment of colours $c : V \rightarrow \{c_1, c_2, c_3\}$ such that every edge of G has endpoints of different colours ($\forall (u, v) \in E, c(u) \neq c(v)$)—special value NIL if this is not possible.

Algorithm: To show that 3-COLOUR-SEARCH is polytime Turing-reducible to 3-COLOUR, assume that algorithm CD(G) solves 3-COLOUR (for every graph G , CD(G) = TRUE iff G is 3-colourable). We write an algorithm CS to solve 3-COLOUR-SEARCH.

CS(G):

if not CD(G): **return** NIL
 # At this point, we know G is 3-colourable.

```

# Assign a first arbitrary vertex to colour 1.
Pick a vertex  $u_1 \in V$ , and let  $C_1 \leftarrow \{u_1\}$ 
# Find every other vertex that can be assigned the same colour as  $u_1$ .
for all  $v \in V - \{u_1\}$ :
    if  $(u_1, v) \notin E$  and  $CD(G_{u_1v})$ :
         $C_1 \leftarrow C_1 \cup \{v\}$  #  $v$  and  $u_1$  can be assigned the same colour
         $G \leftarrow G_{u_1v}$  # continue with  $u_1$  and  $v$  merged, since they have the same colour
# At this point,  $G$  is the result of merging every vertex from  $C_1$  with  $u_1$ .

# Pick an arbitrary remaining vertex and assign it to colour 2.
Pick a vertex  $u_2 \in V - \{u_1\}$ , and let  $C_2 \leftarrow \{u_2\}$ 
# Find every other vertex that can be assigned the same colour as  $u_2$ .
for all  $v \in V - \{u_1, u_2\}$ :
    if  $(u_2, v) \notin E$  and  $CD(G_{u_2v})$ :
         $C_2 \leftarrow C_2 \cup \{v\}$  #  $v$  and  $u_2$  can be assigned the same colour
         $G \leftarrow G_{u_2v}$  # continue with  $u_2$  and  $v$  merged, since they have the same colour
# At this point,  $G$  is the result of merging every vertex from  $C_2$  with  $u_2$ .

# Since  $G$  was 3-colourable, all remaining vertices can be assigned colour 3.
 $C_3 = V - \{u_1, u_2\}$ 

return  $C_1, C_2, C_3$ 

```

Runtime: Algorithm CS makes $\mathcal{O}(n)$ calls to CD (where $n = |V|$), and merges at most two vertices (in time $\mathcal{O}(n+m)$) for each call to CD. The total running time is therefore $\mathcal{O}(n \cdot t(n, m) + n(n+m))$, where $t(n, m)$ is the running time of CD.

Correctness: The algorithm maintains the loop invariant that G is 3-colourable, and assigns the same colour to different vertices only once it has confirmed that this is possible, based on the following claim.

For every graph $G = (V, E)$ and any $u, v \in V$ such that $(u, v) \notin E$, there is a 3-colouring of G where u and v are assigned the same colour iff G_{uv} is 3-colourable.

Proof: If $G = (V, E)$ is a graph and $u, v \in V$ with $(u, v) \notin E$ and if there is a 3-colouring of G in which u and v are assigned the same colour, then the same colouring can be applied to G_{uv} . Any vertices adjacent in G_{uv} are also adjacent in G so will be assigned different colours.

Conversely, if G_{uv} is 3-colourable, then the same colouring can be applied to the vertices of G , using the same colour for u and v . Again, adjacent vertices will be assigned different colours, because they are also adjacent in G_{uv} .

3. (a) Algorithm:


```

 $A \leftarrow \{(i, j) : 1 \leq i \leq m, 1 \leq j \leq n\}$  #  $A$  is the set of every available intersection
 $S \leftarrow \emptyset$  #  $S$  is the current selection of intersections
while  $A \neq \emptyset$ :
    pick  $(i, j) \in A$  with the maximum value of  $a_{i,j}$ 
    # Add  $(i, j)$  to the selection, then remove it (and all adjacent intersections) from  $A$ .
     $S \leftarrow S \cup \{(i, j)\}$ 
     $A \leftarrow A - \{(i, j), (i-1, j), (i+1, j), (i, j-1), (i, j+1)\}$ 
return  $S$ 

```

Counter-Example:

For the input on the right (where we've indicated the number of accidents prevented for each intersection), the algorithm returns the intersections "5" and "0," for a total of 5 accidents prevented. However, picking the other two intersections would prevent a total of 8 accidents instead.



- (b) Note that the input to the problem can be represented as an undirected graph G , and valid selections of intersections are the same as independent sets in G .

Let $A(G)$ be the smallest number of accidents prevented over all independent sets returned by the greedy algorithm, and $M(G)$ be the *maximum* number of accidents prevented over *all* independent sets of G . We prove that for all inputs G , $A(G) \geq M(G)/4$ —and that for at least one input G_0 , $A(G_0) = M(G_0)/4$, showing that the approximation ratio for our algorithm is equal to 4.

Let S be any independent set returned by the algorithm, and T be any independent set with the maximum number of accidents prevented in G . For all $v \in T$, if $v \notin S$, then there is some intersection $v' \in S$ such that $(v', v) \in E$ —otherwise, v could be added to S . Moreover, $a_{v'} \geq a_v$ because when v was removed from G by the algorithm, it was because of some adjacent intersection v' being added to S , which means that at that point, v' had the largest profit among all remaining intersections, including v .

Since no intersection in S has more than 4 neighbours, for all $v \in S$, there are at most 4 intersections $v_1, v_2, v_3, v_4 \in T$ such that $a_v \geq a_{v_1}, a_v \geq a_{v_2}, a_v \geq a_{v_3}, a_v \geq a_{v_4}$. In other words, for all $v \in S$, $4a_v \geq a_{v_1} + a_{v_2} + a_{v_3} + a_{v_4}$ (in the worst case) to “cover” all intersections in T . Hence, $4 \sum_{v \in S} a_v \geq \sum_{u \in T} a_u$, i.e., $A(G) \geq M(G)/4$, as desired.

To show that $A(G)$ can be equal to $M(G)/4$, consider the input below. The algorithm will select the intersections that prevent $(p+1) + 0 + 0 + 0 + 0 = p+1$ accidents while the maximum number of accidents prevented is obtained by picking the intersections $p + p + p + p = 4p$. This is not quite the desired factor of 4, though it can be made arbitrarily close to it by picking p large enough. To get a factor of 4 exactly, simply set the number of accidents prevented for the middle intersection to p . Then, the selection returned by the algorithm is no longer *guaranteed* to be sub-optimal, but it is *possible* that the algorithm will select the middle “intersection” first, and end up with a solution whose value is exactly 1/4 of the optimum. (Note that the size of this input can be made arbitrarily large by repeating the pattern multiple times, separated by rows and columns of intersections with 0 accidents prevented.)

