

Introduction

Course Information Sheet: waiting for last details to be ironed out, will be distributed later this week; go over main points (marking scheme, textbook).

Problem: collect raw petrol from 100 oil platforms.

(For runtime, assume 10^9 ops/sec.)

- a. Use a tanker
 - parameters: cost(A,B) for any two platforms A, B -- could be different from distance (natural obstacles, etc.)
 - best algorithm: over 4×10^{13} years! (more than 2^{100} ops / 10^9 ops/sec / 31536000 sec/year = 40196936841331 years)
- b. Use pipelines
 - parameters: cost(A,B) for any two platforms A, B; no junctions outside platforms
 - best known algorithm: approx. 10 micro-seconds (100^2 ops)
- c. Add constraint: cannot connect more than 4 pipelines at each platform (or any other fixed constant k instead of 4)
 - best known algorithm: approx. 4×10^{13} years
- d. Ease constraint: allow junctions outside platforms (keep cap on maximum number of pipelines allowed at each junction)
 - best known algorithm: still approx. 4×10^{13} years
- e. Back to tanker: relax requirements to approximate smallest total cost, i.e., cost within factor K of best
 - best known algorithm: still approx. 4×10^{13} years, irrespective of value of K
- f. Add constraint: cost satisfies triangle inequality: for any three platforms A, B, C, $\text{cost}(A,C) \leq \text{cost}(A,B) + \text{cost}(B,C)$ (e.g., if cost is directly proportional to distance)
 - best known algorithm for relaxed requirement: approx. 1 second

What's the deal?

- P: class of problems that have polynomial-time (i.e., "efficient") algorithmic solutions
- NP-hard: class of problems for which no efficient algorithm is known (only known algorithms are exponential time)

Vast majority of real-world problems fall into one of these two classes (P or NP-hard). Important to recognize problems in each class and to handle both kinds of problems appropriately.

In this course:

- techniques for writing efficient algorithms for problems in P;
- techniques for deciding whether a problem is in P or NP-hard;
- techniques for handling NP-hard problems.

Background (from CSC263 and its prerequisites):

- Asymptotic notation (big-Oh, Ω , Θ), analysis of runtimes for iterative and recursive algorithms.
- Data structures: queues, stacks, hashing, balanced search trees, priority queues, heaps, union-find/disjoint sets.
- Graphs: definitions, properties, traversal algos (BFS, DFS).
- Induction and other proof techniques, proving correctness of iterative and recursive algorithms.

Greedy Algorithms

"At each step, make the choice that seems best at the time; never change your mind."

Activity Scheduling.

Input: Activities A_1, A_2, \dots, A_n . Each activity A_i consists of positive integer start time s_i and finish time f_i ($s_i < f_i$).
Output: Subset of activities S such that all activities are "compatible" (no two of them overlap in time) and $|S|$ is maximum.

TERMINOLOGY:

- In general, for maximization problem with solution S worth $\text{val}(S)$,
- "maximAL" = nothing can be added to S to increase $\text{val}(S)$;
 - "maximUM" = no solution has larger value.

A. Brute force: consider each subset of activities.
Correctness? Trivial.
Runtime? $\Omega(2^n)$, not practical.

B. Greedy by start time:

```
sort activities s.t.  $s_1 \leq s_2 \leq \dots \leq s_n$ 
S := {} # partial schedule
f := 0 # last finish time of activities in S
for i in [1,2,...,n]:
    if  $f \leq s_i$ : #  $A_i$  is compatible with S
        S := S U { $A_i$ }
        f :=  $f_i$ 
return S
```

Runtime? Sorting is $\Theta(n \log n)$, main loop is $\Theta(n)$.
Total is $\Theta(n \log n)$.

Correctness? Doesn't work. Counter-example:

```
|-----|
|---| |---| |---| ... |---|
```

C. Greedy by duration:

similar to above except sort by nondecreasing duration, i.e.,
 $f_1 - s_1 \leq f_2 - s_2 \leq \dots \leq f_n - s_n$

Correctness? Counter-example:

```
|-----| |-----| |-----| |-----| ... |-----| |-----|
|---| |---| |---| |---|
```

D. Greedy by overlap count:

similar to above except sort from fewest conflicts to most conflicts
("conflict" = overlap with some other activity)

Correctness? Counter-example:

```
|---| |---| |---| |---| ... |---| |---| |---| |---|
|---| |---| |---| |---| |---| |---| |---|
|---| |---| |---| |---| |---| |---| |---|
```

E. Greedy by finish time:

similar to above except sort by nondecreasing finish time, i.e.,
 $f_1 \leq f_2 \leq \dots \leq f_n$

Correctness? No counter-example...

- Intuition: algorithm picks activities that "free up" resources as early

as possible. BUT: intuition for others also made sense...

- How to tell if this works? Will show general technique for proving correctness of greedy algorithms.
- Let S_0, S_1, \dots, S_n = partial solutions constructed by algo. at the end of each iteration.
- Two possibilities:
 - . Prove each S_i is optimal solution to sub-problem.
Works for some problems, but does not generalize well (some problems don't decompose into sub-problems naturally).
 - . Prove each S_i can be "completed" to reach optimal solution.
Can be trickier but generalizes well.
- Say S_i is "promising" if there is some optimal solution OPT that *extends* S_i using only activities from $\{A_{i+1}, \dots, A_n\}$ (i.e., $S_i \subseteq \text{OPT} \subseteq S_i \cup \{A_{i+1}, \dots, A_n\}$).
Note: OPT may not be unique (there may be more than one way to achieve optimal).
- Prove that " S_i is promising" is a loop invariant, by induction in i (number of iterations).
 - . Base case: $S_0 = \{\}$: any optimal solution OPT extends S_0 using only activities from $\{A_1, \dots, A_n\}$.
 - . Ind. Hyp.: Suppose $i \geq 0$ and optimal OPT extends S_i using only activities from $\{A_{i+1}, \dots, A_n\}$.
 - . Ind. Step: To prove: S_{i+1} is promising w.r.t. $\{A_{i+2}, \dots, A_n\}$.
From S_i to S_{i+1} , algo. either rejects or includes A_{i+1} .
 - Case 1: $S_{i+1} = S_i$
This means A_{i+1} not compatible with S_i . Since OPT includes S_i , A_{i+1} is also incompatible with OPT.
Then OPT extends S_{i+1} using only activities from $\{A_{i+2}, \dots, A_n\}$ (since $S_i \subseteq \text{OPT}$ and A_{i+1} not compatible with $S_i \subseteq \text{OPT}$).
 - Case 2: $S_{i+1} = S_i \cup \{A_{i+1}\}$
OPT may or may not include A_{i+1} , so consider both possibilities.
 - Subcase 2.1: $A_{i+1} \in \text{OPT}$
Then OPT already extends S_{i+1} using only activities from $\{A_{i+2}, \dots, A_n\}$.

NOTE: Every case and subcase so far holds no matter how the activities are sorted initially -- in other words, our proof does not yet depend on the ordering. But we know this is important: it comes into the next subcase.

- Subcase 2.2: $A_{i+1} \notin \text{OPT}$
How can this happen? There must be $A_j \in \text{OPT}$ that overlaps with A_{i+1} (otherwise, $\text{OPT} \cup A_{i+1}$ would be better than optimal OPT). Also, $j > i+1$ because A_{i+1} is compatible with S_i , so $f_j \geq f_{i+1}$ and at most one A_j overlaps A_{i+1} (otherwise OPT would contain overlapping activities, or an activity outside S_i with finish time earlier than

A_{i+1} , both impossible).
But then, $OPT' = OPT \cup \{A_{i+1}\} - \{A_j\}$ extends S_{i+1}
using $\{A_{i+2}, \dots, A_n\}$: same number of activities as OPT ,
and no overlap introduced because $f_{i+1} \leq f_j$.

NOTE: Argument above known as "exchange lemma": arguing that any optimal solution can be made to agree with greedy solution, one element at a time. Just like definition of "extends", every problem and algorithm yields a different "exchange lemma" -- there is no single Exchange Lemma that applies to every algorithm and problem!

In all cases, there is some optimal OPT' that extends S_{i+1} using only activities from $\{A_{i+2}, \dots, A_n\}$.

- So each S_i is promising. In particular, S_n is promising, i.e., there is optimal OPT that "extends" S_n using only activities from $\{\}$. In other words, S_n is optimal.

For Next Week

- * Readings: Sections 5.1, 4.4.
- * Self-Test: Exercises 5.1, 5.2, 4.1.