

Minimum Spanning Tree (Review from CSC263)

Input: Connected undirected "weighted" graph $G = (V, E)$, i.e., with non-negative integer weight/cost $w(e)$ for each edge $e \in E$.

Output: A spanning tree $T \subseteq E$ such that $\text{cost}(T)$ (sum of the costs of edges in T) is minimal.

- Terminology:

- . "Spanning tree": acyclic connected subset of edges.
- . "Acyclic": does not contain any cycle.
- . "Connected": contains a path between any two vertices.

- Properties: For all spanning trees T of a graph G ,

- . T contains exactly $n-1$ edges;
- . the removal of any edge from T disconnects it into two disjoint subtrees (otherwise T would contain a cycle);
- . the addition of any edge to T creates exactly one cycle (otherwise T would be disconnected).

A. Brute force: consider each possible subset of edges.

Runtime? Exponential, even if we limit search to spanning trees of G (instead of considering all possible subsets of edges).

B. Prim's algorithm:

Idea: Start with some vertex $r \in V$ (pick arbitrarily) and at each step, add lowest-cost edge that connects a new vertex to existing partial tree.

Runtime? $\Theta(m \log n)$ using min-heap to implement priority queue of candidate edges (with priority = weight of smallest edge to tree).

C. Kruskal's algorithm:

Idea: Repeatedly put in smallest-cost edge remaining, as long as it doesn't create a cycle.

Runtime? $\Theta(m \log m)$ for sorting; main loop involves sequence of m Union and FindSet operations on n elements which is $\Theta(m \log n)$ -- faster using best heuristics. Total is $\Theta(m \log n)$ since $\log m$ is $\Theta(\log n)$.

D. Reverse-delete algorithm:

Idea: Repeatedly delete largest-cost edge remaining, as long as it does not disconnect the graph.

Runtime? For each edge, run BFS/DFS starting from one endpoint to figure out if other endpoint can still be reached. Requires $\Omega(m^2)$.

Correctness of Kruskal's algorithm:

- First, spell out algorithm in pseudo-code:

```
Sort edges by non-decreasing weight:  $w(e_1) \leq \dots \leq w(e_m)$ .
 $T = \{\}$ 
for  $j = 1, 2, \dots, m$ :
    let  $(u, v) = e_j$ 
    if  $T$  does not contain a path between  $u$  and  $v$ :
         $T = T \cup \{e_j\}$ 
```

- Algorithm generates subsets of edges T_0, T_1, \dots, T_m .

Say T_i is "promising" if it can be extended to some MST T^* using only edges $\{e_{i+1}, \dots, e_m\}$, i.e., $T_i \subseteq T^* \subseteq T_i \cup \{e_{i+1}, \dots, e_m\}$.

(Set notation again because solution is a subset of input again.)

- Loop invariant: T_i is promising.

Base: $T_0 = \{\}$ is promising: every MST T^* is a subset of $\{e_1, \dots, e_m\}$.

I.H.: Suppose $i \geq 0$ and T_i can be extended to T^* .

Step: Either $T_{i+1} = T_i$ or $T_{i+1} = T_i \cup \{e_{i+1}\}$.

Case 1: If $T_{i+1} = T_i$, then $T_i \cup \{e_{i+1}\}$ contains a cycle.
Since $T_i \subset T^*$ and T^* is a MST, $e_{i+1} \notin T^*$, so
 $T_{i+1} \subset T^* \subset T_{i+1} \cup \{e_{i+2}, \dots, e_m\}$, i.e.,
 T^* extends T_{i+1} .

Case 2: If $T_{i+1} = T_i \cup \{e_{i+1}\}$, then consider whether or not
 $e_{i+1} \in T^*$.

Subcase 2.1: If $e_{i+1} \in T^*$, then
 $T_{i+1} \subset T^* \subset T_{i+1} \cup \{e_{i+2}, \dots, e_m\}$, i.e.,
 T^* extends T_{i+1} .

Subcase 2.2: If $e_{i+1} \notin T^*$, then T^* does not extend
 T_{i+1} . Construct T^{**} that does, as follows.
Consider endpoints of e_{i+1} in T^* : they are connected by a
path. Fact: not all edges on this path belong to T_i -- else
algorithm would not generate $T_{i+1} = T_i \cup \{e_{i+1}\}$. So
 T^* contains some edge e_j on this path with $j > i+1$ (because
 T^* agrees with T_i on all edges e_1, \dots, e_i and e_{i+1} is
not in T^*).
Then $w(e_j) \geq w(e_{i+1})$ and we let
 $T^{**} = T^* \cup \{e_{i+1}\} - \{e_j\}$.
 T^{**} is a MST: it is connected, acyclic, and with total cost
 \leq total cost of T^* -- in fact, it must be that $w(e_j) =$
 $w(e_{i+1})$ since T^* is also optimal.

In every case, T_{i+1} is promising.

Since every T_i is promising, T_m is promising: $T_m = T^*$ for some MST
 T^* . So T_m is a MST.

- Correctness of other algorithms proved similarly.

Shortest Paths

Input: connected graph $G = (V, E)$ with edge costs $w(e)$ for all $e \in E$;
vertices $s, t \in V$. IMPORTANT: costs are POSITIVE integers.

Output: a path from s to t with minimum total cost ("shortest" path).

- Brute-force: in general, exponentially many paths possible.

- Special case: if $w(e) = 1$ for all e : BFS!

- Dijkstra's algorithm: "modified" BFS: use `_priority queue_` instead of
queue to collect unvisited vertices; set `priority` = shortest distance so
far. Note similarity to Prim's algorithm, but also important difference:
for Prim's algorithm, `priority` = minimum cost of single edge to new
vertex; for Dijkstra's algorithm, `priority` = minimum total distance from
 s to new vertex.

Intuition: find shortest s - t path by finding shortest paths from s to
every vertex.

- Algorithm:

```
# Initialization.
P = {} # edges in shortest paths tree
initialize empty min-priority queue
for all v in V:
    pi[v] = NIL # predecessor of v on shortest s-v path so far
    d[v] = oo # priority of v = minimum distance s-v so far
    enqueue v # with priority d[v] = oo
d[s] = 0
update queue order of s

# Main loop.
while queue not empty:
    v = dequeue element with minimum priority d[]
    P = P u {(pi[v],v)} # problem when v = s: handled below
    for all edges (v,u):
        if u in the queue and d[v] + w(v,u) < d[u]: # "relaxation"
            pi[u] = v
            d[u] = d[v] + w(v,u)
            update queue order of u
P = P - {(NIL,s)} # clean up

return P
```

- Runtime:

- $O(n)$ for initialization.
- Main loop iterates at most n times (each iteration removes one vertex from the queue).
- Each iteration examines edges in one adjacency list and updates priorities. Over all iterations, each edge generates at most one queue update. And each priority update takes time $O(\log n)$.
- Total: $O(m \log n)$.

- Optimal substructure in Shortest Path problem:

A subpath of a shortest path is also a shortest path.

Proof: Cut and paste technique! Assume that we have a shortest path (P) between two vertices u and v and consider two vertices x and y on this path. If the path between x and y along P is not a shortest path between x and y , we can simply replace this subpath by the shortest between x and y (and we are sure that there is one, why?). Replacing this subpath with another shorter path will make the shortest path between u and v even smaller which is in contradiction with the fact that path p is a shortest path! The reasoning applies even if the shortest x - y path contains nodes that appear elsewhere on P -- then, the contradiction is derived from using only the non-overlapping parts of the path.

- Triangle Inequality:

for all vertices u, v, x in V we have
 $\delta(u,v) \leq \delta(u,x) + \delta(x,v)$.

- Lemma:

During Dijkstra's algorithm, $d[v] \geq \delta(s,v)$ and this holds after taking any sequence of relaxations.

Proof idea: by induction. The inequality holds after the initialization step (every vertex except s has $d[v] = \text{infinity}$ and $d[s] = 0$). And each relaxation preserves the property (because $d[v]$ is either ∞ or equal to the total weight of some path in G).

- Correctness (main idea):

Show $d[v] = \delta(s,v)$ for all v in P_i , where $\delta(x,y)$ is minimum total weight of any x - y path.

Consider one iteration of main loop: v with minimum $d[v]$ removed from queue, $P_{i+1} = P_i \cup \{(p_i[v], v)\}$.

To show: $d[v] = \delta(s,v)$. For contradiction, suppose $d[v] > \delta(s,v)$ and let $P = s$ - v path with total weight $\delta(s,v)$. Let (x,y) be first edge on P with x in P_i , y outside P_i . Either $y = v$ or $y \neq v$.

Case 1: $y = v$.

Then $d[v] \leq d[x] + w(x,v)$ [(x,v) one possible edge from P_i to v]
 $= \delta(s,x) + w(x,v)$ [by I.H. since x in P_i]
 $= \delta(s,v)$
 $< d[v]$,

contradiction: $d[v] < d[v]$!

Case 2: $y \neq v$.

Then $d[y] \leq d[x] + w(x,y)$ [(x,y) one possible edge from P_i to y]
 $= \delta(s,x) + w(x,y)$ [by I.H. since x in P_i]
 $< \delta(s,x) + w(x,y) + \delta(y,v)$ [positive weights]
 $= \delta(s,v)$
 $< d[v]$,

contradiction: $d[y] < d[v]$ but $d[v] = \text{minimum outside } P_i$.

So $d[v] = \delta(s,v)$.

Rest of proof involves induction structure around main idea. Included below for reference but not covered during lecture.

- Correctness (details):

Algorithm generates subsets of edges P_0, P_1, \dots, P_{n-1} .

Say P_k is "promising" if it can be "extended" to some collection of shortest paths P^* (really, a shortest paths tree) using only edges that do not have *both* endpoints "in" P_k , i.e., edges with at least one endpoint still in the queue. (Technically, P_k contains edges, not vertices: when we speak of a vertex being "in" P_k , we mean that it is the endpoint of some edge in P_k .)

Loop invariant:

- P_k is promising, and
- for all u in P_k , all v outside P_k ,
 $d[u] = \delta(s,u) \leq \delta(s,v) \leq d[v]$

where $\delta(s,v)$ is the minimum total weight of all paths from s to v . (Extra clause is required for the proof to go through.)

Proven by induction on k .

Base Case:

$P_0 = \{\}$ is promising, trivially.

Ind. Hyp.:

For some arbitrary k , suppose P_k can be extended to some shortest paths tree P^* , using only edges without both endpoints in P_k , and that $d[u] = \delta(s,u) \leq \delta(s,v) \leq d[v]$ for all u in P_k and v outside P_k .

Ind. Step:

Consider $P_{k+1} = P_k \cup \{(u,v)\}$, with u in P_k and v outside P_k .

Either $(u,v) \in P^*$ or it does not.

If $(u,v) \in P^*$, then P^* extends P_{k+1} .

Also, $\delta(s,v) = \delta(s,u) + w(u,v)$ (because $(u,v) \in P^*$) so $d[v] = d[u] + w(u,v) = \delta(s,u) + w(u,v) = \delta(s,v)$.

Moreover, because $d[v]$ was the smallest of the $d[]$ values for vertices outside P_k , $d[x] = \delta(s,x) \leq \delta(s,w) \leq d[w]$ for all x in P_{k+1} and all w outside P_{k+1} .

If $(u,v) \notin P^*$, then:

- consider the path P in P^* from s to v , and let (w,v) be the last edge on this path;
- w must belong to P_k -- otherwise, let (x,y) be the first edge on P with x in P_k , y outside P_k
 - . $d[y] \leq d[x] + w(x,y)$ (because $d[y]$ is the smallest value of $d[t] + w(t,y)$ for all edges (t,y) with t in P_k)
 - $= \delta(s,x) + w(x,y)$ (because x in P_k)
 - $< \delta(s,x) + w(x,y) + \delta(y,v)$
 - (since all edge weights strictly positive)
 - $= \delta(s,v)$
 - $\leq d[v]$
 - . but this would contradict the fact that $d[v]$ is the smallest value of $d[t]$ for all vertices t outside P_k
- so w in P_k ;
- so $\delta(s,v) = \delta(s,w) + w(w,v) = d[w] + w(w,v)$;
- since $d[v]$ is the minimum of $d[x] + w(x,v)$ over vertices x , this means $d[v] \leq d[w] + w(w,v) = \delta(s,v)$;
- so $d[v] = \delta(s,v)$ (it cannot be smaller);
- so we can let $P^{**} = P^* - \{(w,v)\} \cup \{(u,v)\}$, and after the update to $d[w]$ for all edges (v,w) with w outside P_k , we still have that $d[x] = \delta(s,x) \leq \delta(s,y) \leq d[y]$ for all x in P_k , y outside P_k .

Hence, the loop invariant holds. When the algorithm terminates, this means $d[u] = \delta(s,u)$ for all vertices u : we have found shortest paths to every vertex.

For Next Week

- * Readings: Section 6.5.
- * Self-Test: Trace matrix chain multiplication algo. on an example.