

Choosing augmenting paths efficiently:

- * Example where it could take $2 * 10^{100}$ augmentations to find max flow of $2 * 10^{100}$. In worst-case, Ford-Fulkerson takes time $\Theta(m f^*)$, where f^* is the value of a max flow -- this is not polytime.
- * Edmonds-Karp algorithm: use BFS on residual network to find augmenting paths; guaranteed to find an augmenting path with smallest # of edges in time $O(m)$. Possible to prove that no more than $O(mn)$ augmentations are required to find max. flow. Total time: $O(nm^2) = O(n^5)$.
- * Dinic algorithm: perform complete BFS, use all augmenting paths w.r.t. BFS tree, then repeat. Worst-case time down to $O(n^2m) = O(n^4)$.
- * "Preflow" algorithms: don't use augmenting paths; instead, push as much as possible along individual edges then go back to fix conservation. More complicated to explain and write down correctly, but cuts down time to $O(n^3)$.
- * Why are we not concerned about details? Applications of network flow do not involve writing new algorithms: you always solve the same problem, so you can use an existing implementation that's already been debugged and optimized. Applications involve taking a problem, casting it in the guise of a network flow, solving the network flow problem, then casting the answer back to your problem.

Applications of network flows:

- * Maximum bipartite matching:
 Input: An undirected bipartite graph $G=(V_1, V_2, E)$ -- one where every edge is between V_1 and V_2 (i.e., no edge has both endpoints in the same "side").
 Output: A disjoint subset of edges (a "matching") with maximum size ("disjoint" edges means no edge shares an endpoint with any other edge in the matching).

Algorithm:

1. Create network N from G : add source s , sink t , edges from s to every vertex in V_1 , edges from every vertex in V_2 to t ; make original edges directed from V_1 to V_2 ; set all edge capacities to 1.
2. Find maximum flow in N .
3. Output $M = \{\text{every original edge with flow} = 1\}$.

Correctness?

- Any matching in graph yields flow in network: set flow = 1 for graph edges in matching, 0 for graph edges not in matching; set flow = 1 for new edges to/from matched vertices, 0 for new edges to/from unmatched vertices.
- Any integer flow in network yields matching in graph: pick edges with flow = 1 (leave out edges with flow = 0). No edge can have flow larger than 1 because of capacities.

For this correspondence, size of matching = value of flow. Hence, any max flow in network yields a maximum matching in graph (because a larger matching would give a larger flow).

Difference between Network Flow and other techniques:

* When solving optimization problem with "network flow techniques", what we are doing is actually a **reduction** or **transformation**: we take the input to our problem and create a network from it, then use standard algorithm to solve the maximum flow problem on that network (it's always the same problem and always the same algorithm, only the input differs). Then, we use the solution to the network flow problem to reconstruct a solution to our problem.

So the algorithm we end up with will always have the following structure

1. Describe how to create network N from input x (to original problem)
2. Find max flow f (or min cut X) in N , using standard algorithms
3. Describe how to construct solution to x from f (or X)

There are two ways this could go wrong:

- The network we construct could fail to represent certain solutions to our original problem. We show this is **not** the case by arguing that every solution to the original problem yields a valid flow (or cut) in the network.
- The network we construct could have solutions that don't correspond to anything in our original problem. We show this is **not** the case by arguing that every flow (or, depending on the problem, every **integer** flow, or every cut) in the network corresponds to a solution to the original problem.

Sometimes those arguments can be very short, because the correspondence is obvious between both problems. But both arguments are important and must always be included.

Another application: Maximum Independent Set in a bipartite graph

Input: Bipartite graph $G = (V_1, V_2, E)$.

Output: Independent Set $I \subseteq V_1 \cup V_2$ with $|I|$ maximum -- vertex subset I is "independent" iff no edge has both endpoints in I .

Algorithm:

1. Create network as for bipartite matching except $c(u, v) = \infty$ for each edge (u, v) in original graph.
2. Find cut $X = (V_s, V_t)$ with minimum capacity $c(X)$ -- first, find max flow, then run algo from proof of Ford-Fulkerson theorem to generate min-capacity cut.
3. Return $I = (V_1 \cap V_s) \cup (V_2 \cap V_t)$.

Correctness?

- For any independent set I , let $V_s = \{s\} \cup (V_1 \cap I) \cup (V_2 - I)$; $V_t = V - V_s$. Then the only edges across the cut $X = (V_s, V_t)$ are those from s to $V_1 - I$ and those from $V_2 - I$ to t . So $c(X) = |V_1| + |V_2| - |I|$. So min cut capacity is at most $|V_1| + |V_2| - |I|$, where I is max independent set.
- For any cut $X = (V_s, V_t)$ with finite capacity, let $I = (V_1 \cap V_s) \cup (V_2 \cap V_t)$. Because $c(X)$ is finite, there is no edge (u, v) from original graph with u in V_s , v in V_t , so I is an independent set. Moreover, the only edges across X are those from s to $V_1 - I$, and those from $V_2 - I$ to t , which means $c(X) = |V_1| + |V_2| - |I|$. So max independent set has size at least $|V_1| + |V_2| - c(X)$, where X is min-capacity cut.

Hence, max independent set size = min cut capacity and the algorithm correctly returns an independent set of maximum size.

For Next Week

- * Readings: Sections 7.1
- * Self-Test: Exercises 7.17(a), 7.17(b)