

Coping with NP-completeness

So far...

- * Techniques for writing (hopefully) efficient algorithms (greedy, dynamic programming, network flows, linear programming).
- * Techniques for showing that problems cannot be solved efficiently (if $P \neq NP$, then every NP-complete problem D does not belong to P).
- * Traditional point of view:
 - P = "easy"
 - NP-complete = "hard"
- * But:
 - definition of NP-completeness based on worst-case analysis (maybe inputs encountered in practice are rarely worst-case)
 - for "real-world" input sizes ($\geq 10^9$), even n^2 runtime is inefficient!

Approximation algorithms

Not possible to solve NP-hard problems *exactly* and *in polytime* (unless $P = NP$).

In practice, sometimes sacrifice on efficiency: use exponential-time algorithm and hope inputs don't trigger worst-case behaviour. Particularly useful on restricted families of inputs:

- * A problem that's hard in general may be easy for inputs of a certain type. For example:
 - 2SAT.
 - Independent set on a tree.
 - Many graph problems are easy on trees or other restricted kinds of graphs. Others not (e.g., planar 3-colouring still NP-complete).

Sometimes sacrifice on exactness, particularly for optimization problems: instead of searching for *best* solution, settle for "good enough" solution. But what does that mean exactly?

For minimization problems, let $OPT(x)$ be the minimum *value* of any solution for input x . Suppose we have an approximation algorithm that generates solutions with approximate value $A(x)$. By definition, $OPT(x) \leq A(x)$ for all inputs (since $OPT(x)$ is minimum).

The "approximation ratio" of our algorithm is a function $r(n)$ such that $A(x) \leq r(n) * OPT(x)$ for all n and all inputs x of size n , i.e., approximation ratio gives a bound on how much larger than optimum our approximate value might be -- it gives a guarantee that approximate values cannot be "too" large compared to optimum.

Vertex Cover:

- * Approx algo 1: greedy strategy -- next week's tutorial...

* Approx algo 2:

Repeatedly pick an edge and put both endpoints in C, then remove all edges incident on the two endpoints, until no edge remains. Then,

$$|C| \leq 2 * OPT$$

because all covers include at least one endpoint from every edge in C (all edges in C are disjoint, i.e., with no endpoint in common) so in particular, $OPT \geq |C|/2$. This shows approximation ratio ≤ 2 .

To show approximation ratio = 2, need an example where algorithm performs that badly -- in this case, use n disjoint edges! Algorithm returns 2n endpoints but n of them are sufficient.

Q: In general, how can we compute ratio without knowing OPT?

(Particularly for NP-hard problems, like VC, TSP, bin packing?)

A: Use a lower bound. Find another value LB that's easy to compute and for which you can prove $LB \leq OPT$ and $A \leq r * LB$. For example...

* Approx algo 3:

- Create linear program from input graph:

variables: x_1, \dots, x_n (one for each v_i in V)

obj. function: minimize $x_1 + \dots + x_n$

constraints: $0 \leq x_i \leq 1$ for all i

$x_i + x_j \geq 1$ for each (v_i, v_j) in E

(As an Integer Program -- domain for each $x_i = \mathbb{Z}$ -- this is completely equivalent to MinVertexCover, including NP-hardness.)

- Compute optimal solution to linear program: $x^*_1, x^*_2, \dots, x^*_n$ (linear program "relaxation": allow real values for variables).

- Create cover as follows:

for each v_i in V, put v_i in C iff $x^*_i \geq 1/2$.

(C is a cover because constraint $x_i + x_j \geq 1$ guarantees at least one of $x^*_i, x^*_j \geq 1/2$ for each edge (v_i, v_j) .)

Approximation ratio?

Consider minimum vertex cover C' . For $i = 1, \dots, n$, let $x'_i = 1$ if v_i in C' ; $x'_i = 0$ otherwise. x'_1, \dots, x'_n is a solution to linear program that satisfies all constraints with 0-1 values so

$$|C'| = \sum x'_i \geq \sum x^*_i$$

where x^*_i is optimal solution to linear program with no restriction on values, so guaranteed to be at least as small as any other solution, including those with additional restrictions.

For $i = 1, \dots, n$, let $\tilde{x}_i = 1$ if $x^*_i \geq 1/2$; $\tilde{x}_i = 0$ otherwise.

Then, for each i, $\tilde{x}_i \leq 2 x^*_i$ so

$$|C| = \sum \tilde{x}_i \leq 2 \sum x^*_i \leq 2 |C'| \quad (\text{by equation above})$$

Hence, $|C|$ is no more than twice the size of a minimum vertex cover.

How well can problems be approximated?

Even though all NP-complete problems "equivalent" to each other (in one sense), approximation ratios for corresponding optimization problems all over the place.

* VC: constant approx ratio 2.

* Set Cover (in textbook): approx ratio $\log n$ (not constant but limited).

* Knapsack: approx ratio $1+\epsilon$ in time $O(n^3/\epsilon)$, for all constants ϵ in $(0,1]$!

* TSP: no finite ratio, unless $P=NP$!

Traveling Salesman Problem (TSP):

* Given graph G with edge weights $w(e)$, find a "tour" (Hamiltonian cycle) with minimum total weight.

* NP-hard: no polytime solution.

* NP-hard to approximate with constant ratio:

Suppose we have an algorithm with approx ratio $C(n)$, i.e., guaranteed to find tour with total weight $\leq C * \text{OPT}$. We show how this algorithm could be used to solve an NP-hard problem.

Given an input G to HamCycle problem, construct an instance of TSP G' as follows: on the same vertex set as G , put in all possible edges with $w(e) = 1$ if e in G , $w(e) = Cn+1$ if e not in G .

Any tour in G' that uses only edges with weight 1 has total weight n ;

any tour in G' that uses at least one edge with weight $Cn+1$ has total weight $> Cn$.

If G contains a Ham cycle, then G' contains a tour with total weight n ;

if G does not contain a Ham cycle, then all tours in G' have total weight $> Cn+1$ (they must contain at least one edge not in G).

So now, run approx. algorithm on G' . If algo returns answer with total weight $\leq C*n$, then G contains a Ham cycle; if algo returns answer with total weight $> Cn+1$, then G does not contain a Ham cycle. As this solves the Ham cycle problem, approx algorithm cannot run in polytime.

For Next Week

* Readings: Sections 9.2.3, 9.2.4, 9.2.5.

* Self-Test: None for this week...